

V: Polygonal Meshes & Rendering

3D CV

Kirill Struminsky

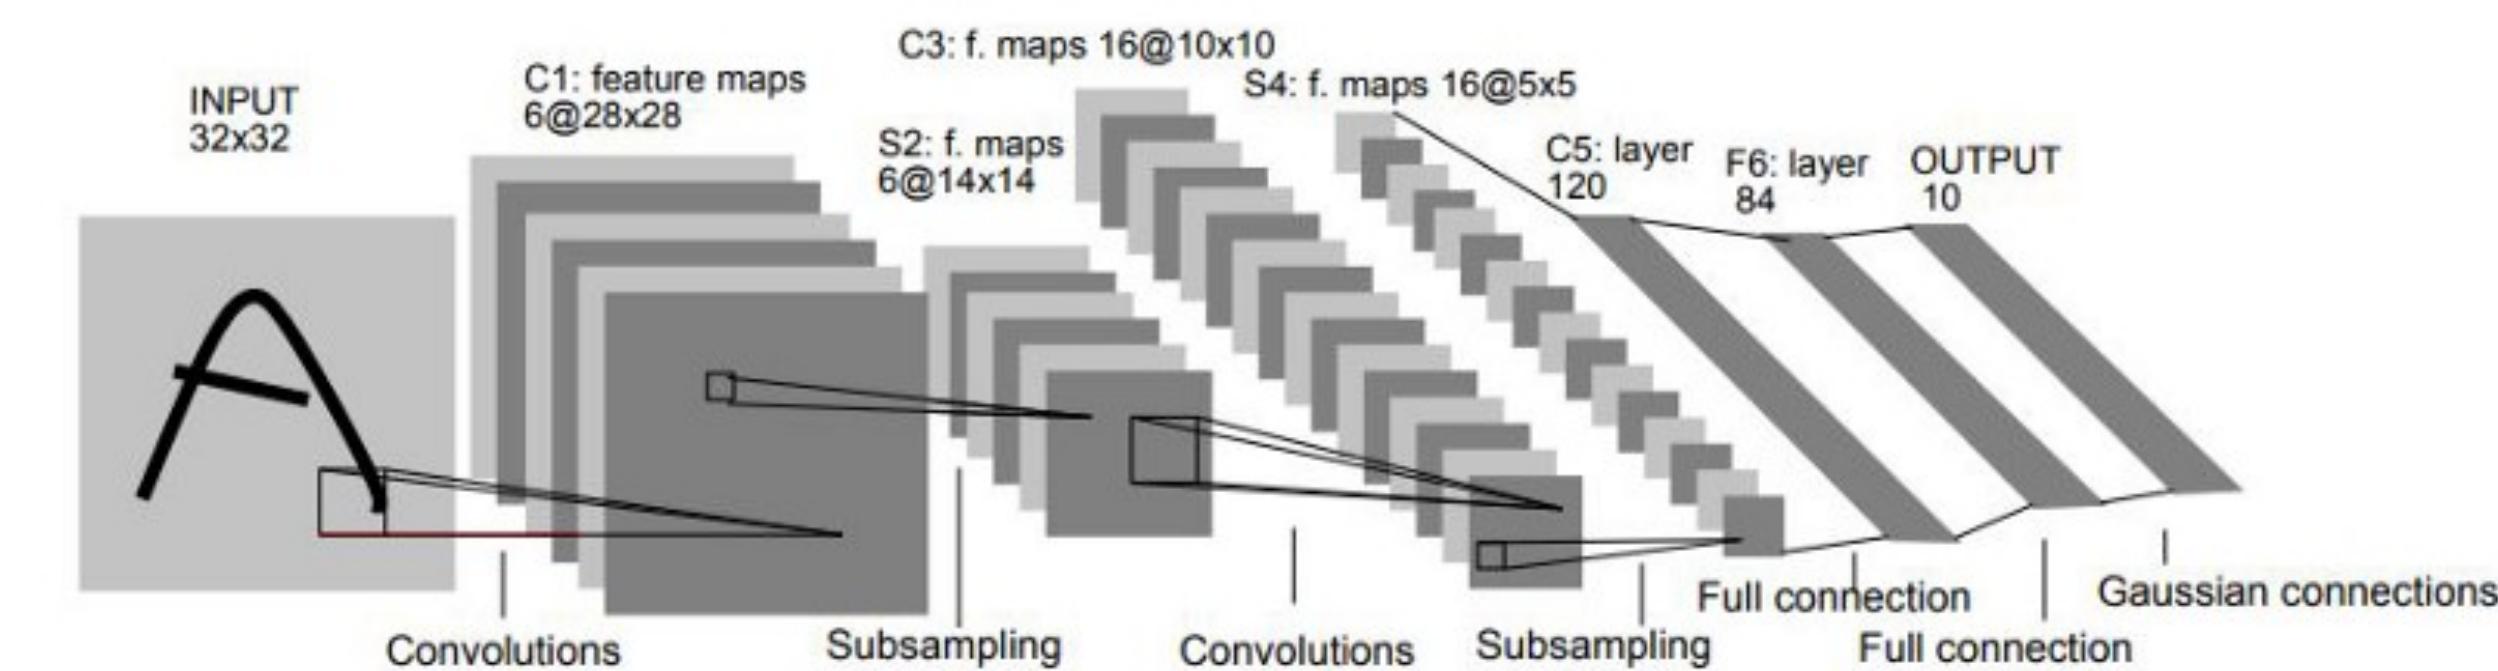
In the Previous Episode

- Point cloud processing
 - Symmetries in deep learning
 - PointNet and it's successors
- Today: polygonal meshes, going beyond point clouds

Data Symmetries

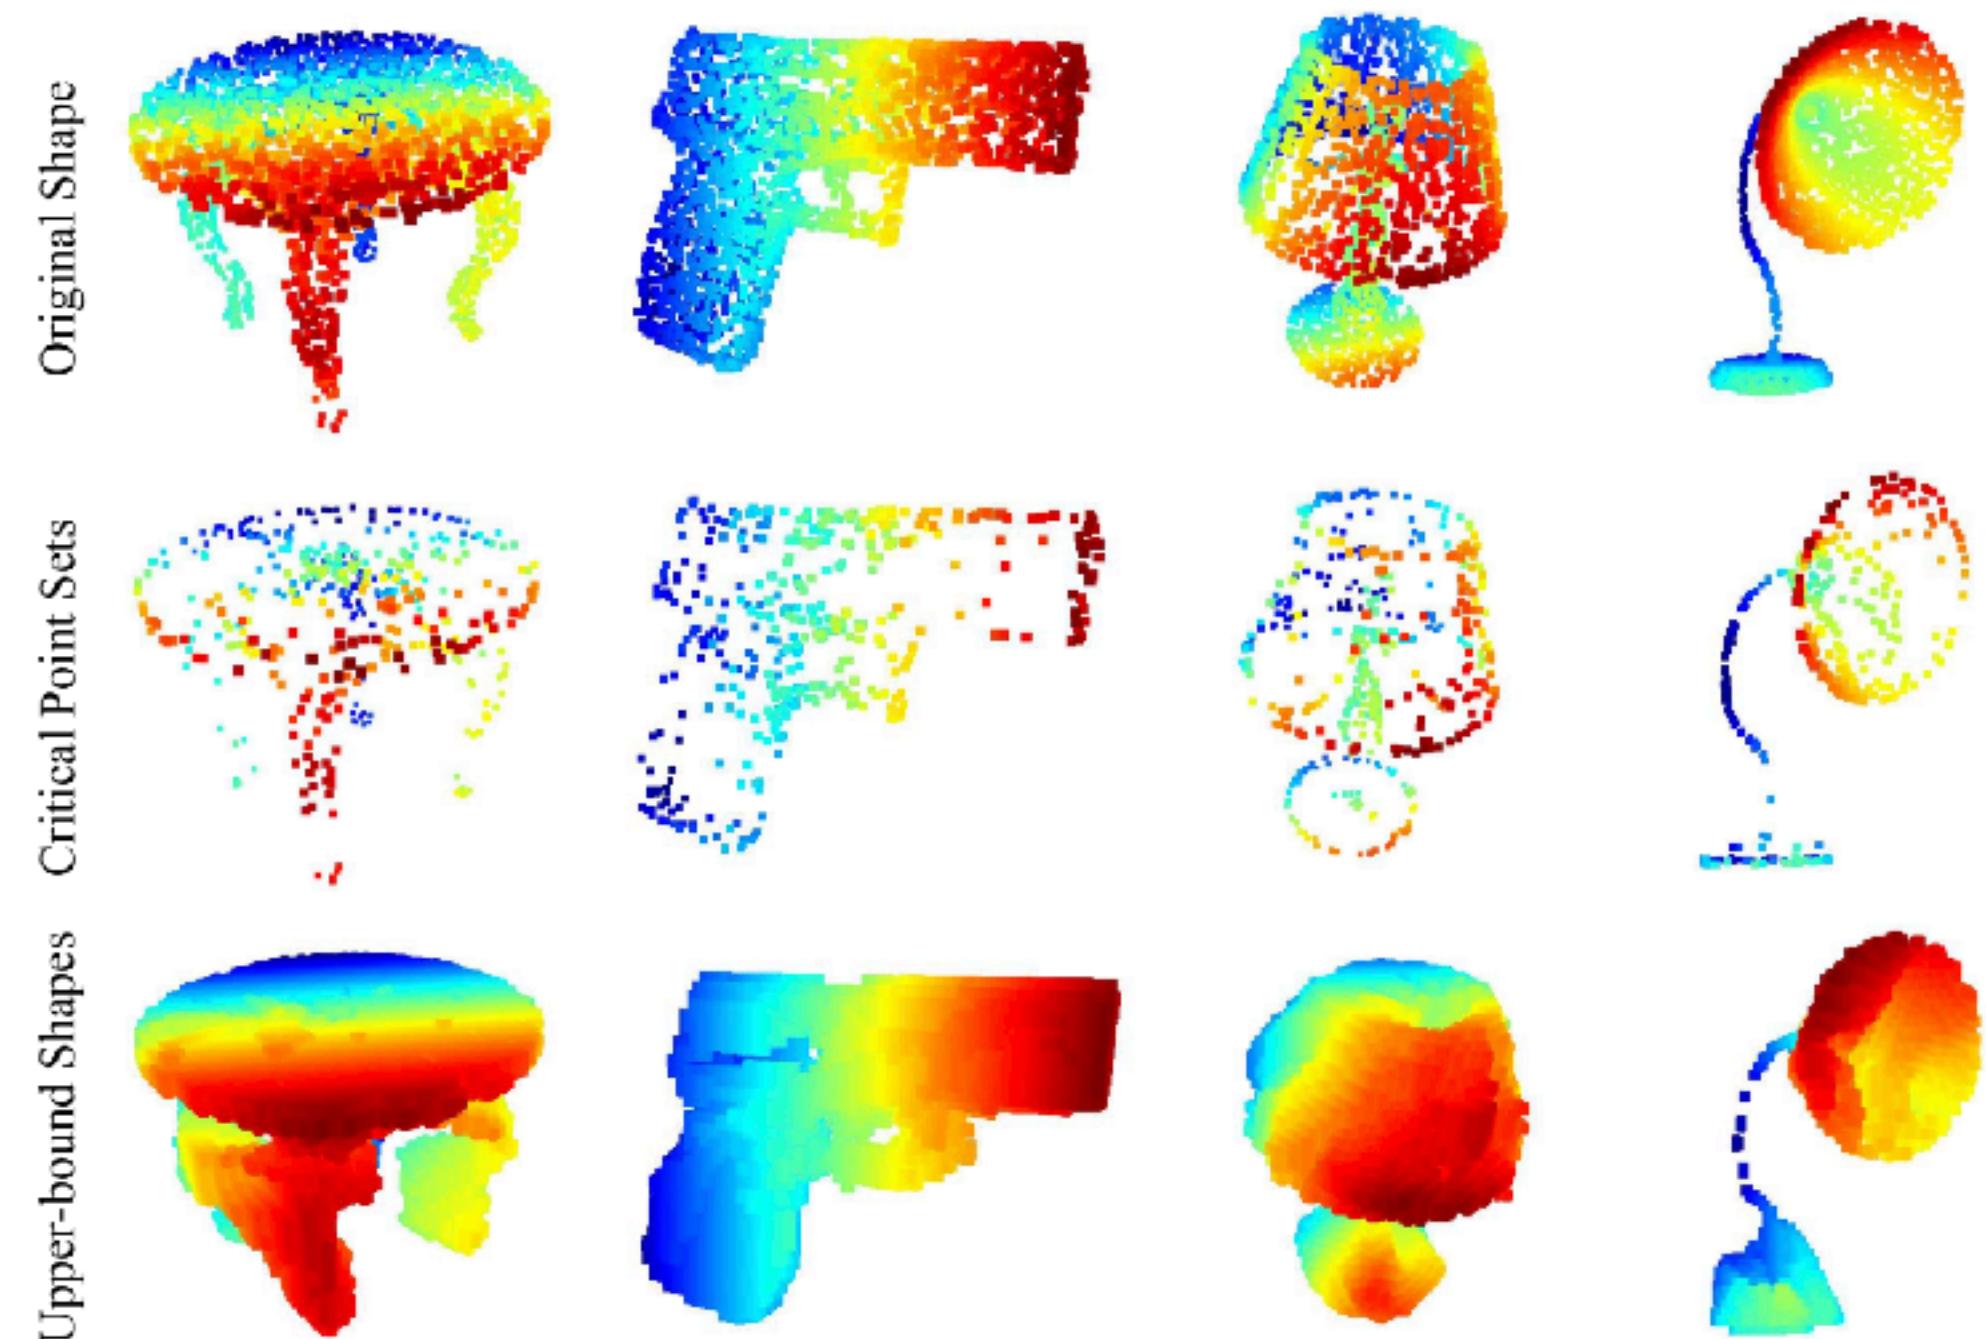
In general

- Consider function f and a group G acting on $\text{dom } f$
- **Equivariance:** $\forall g \in G \ f(g \cdot x) = g \cdot f(x)$
- **Invariance:** $\forall g \in G \ f(g \cdot x) = f(g)$
- Deep architectures exploit symmetries by stacking equivariant layers

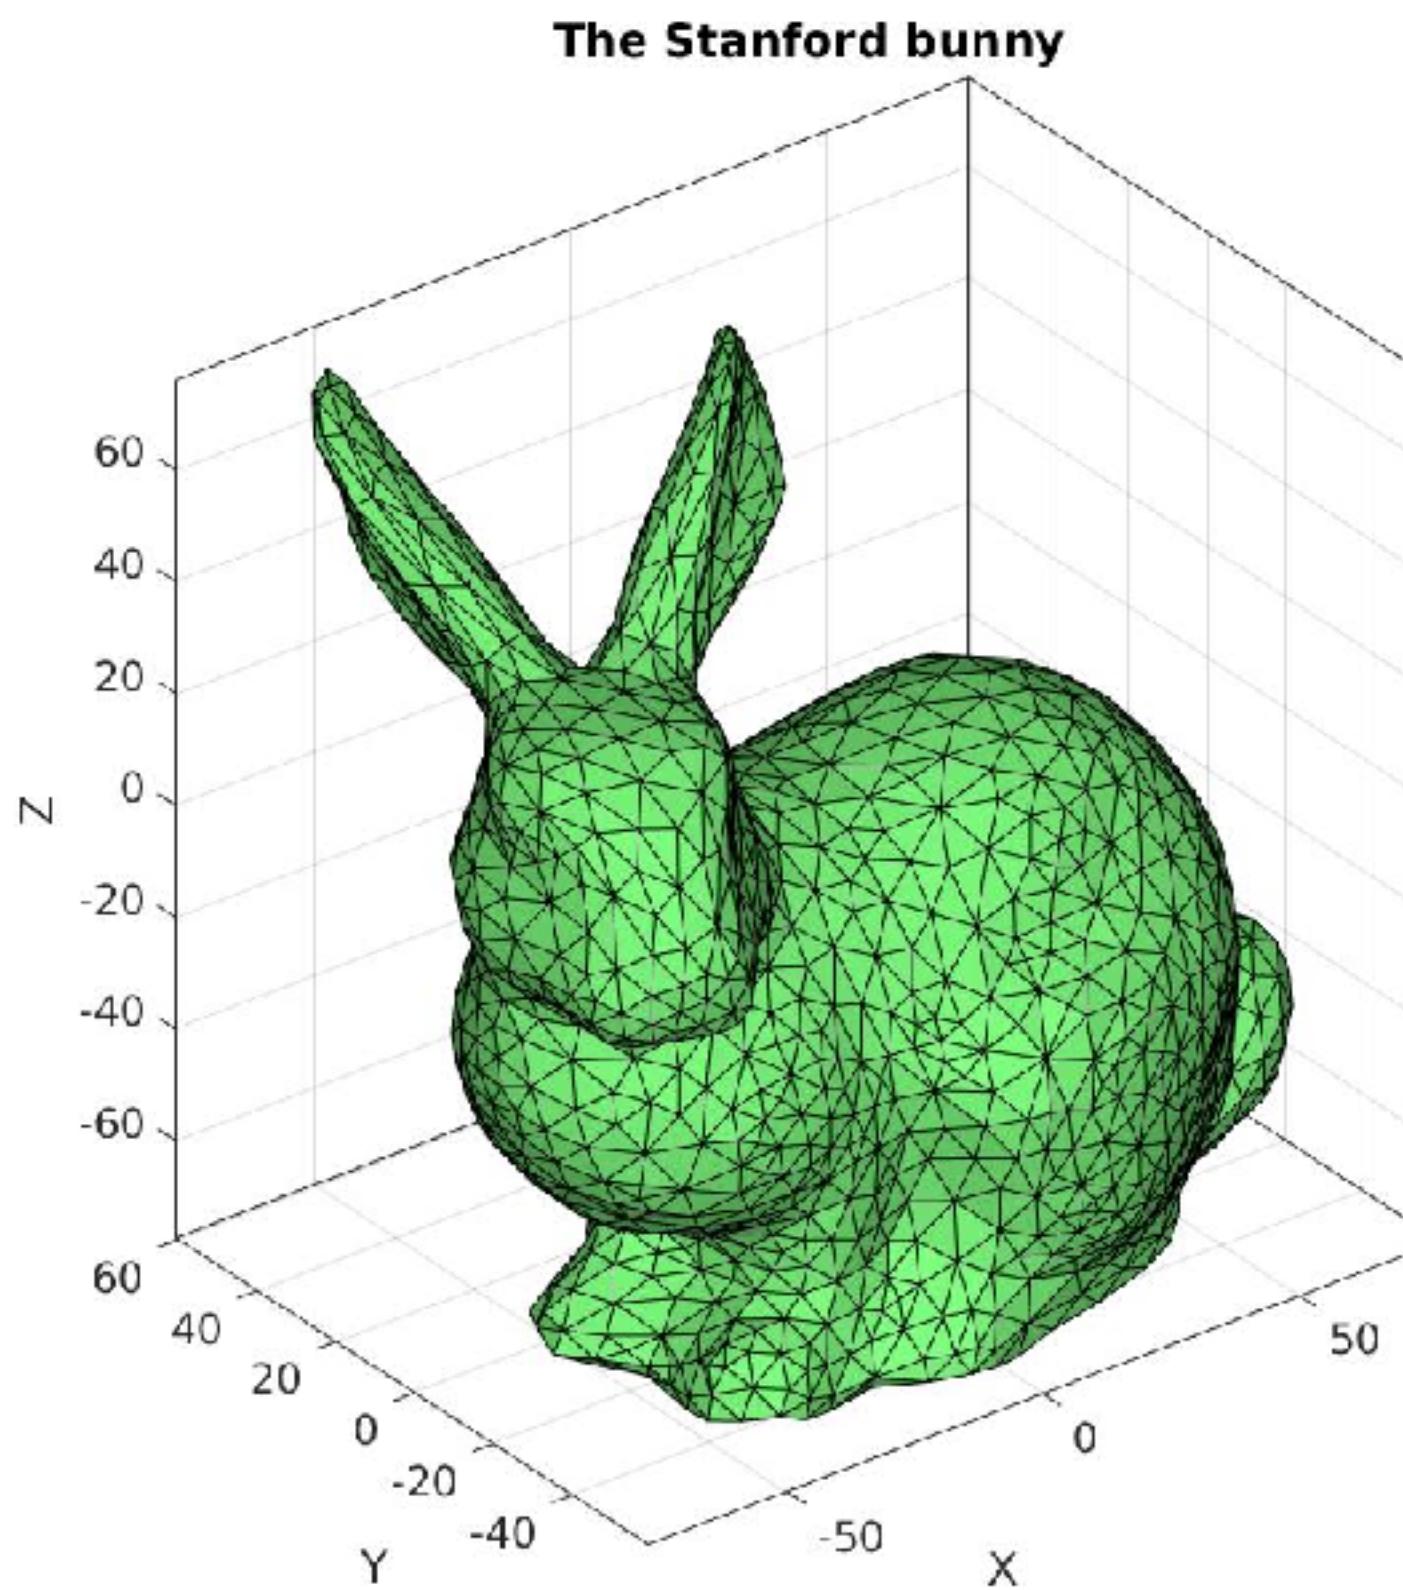


PointNet

- Permutation invariance for learning on sets
 - $f(x_1, \dots, x_n) = g(\text{MAX}(h(x_1), \dots h(x_n)))$
 - In practice, h and g are MLPs
- Rotation invariance with Vector Neurons
- Hierarchical features & Transformers



Today: Meshes & (Differentiable) Rendering



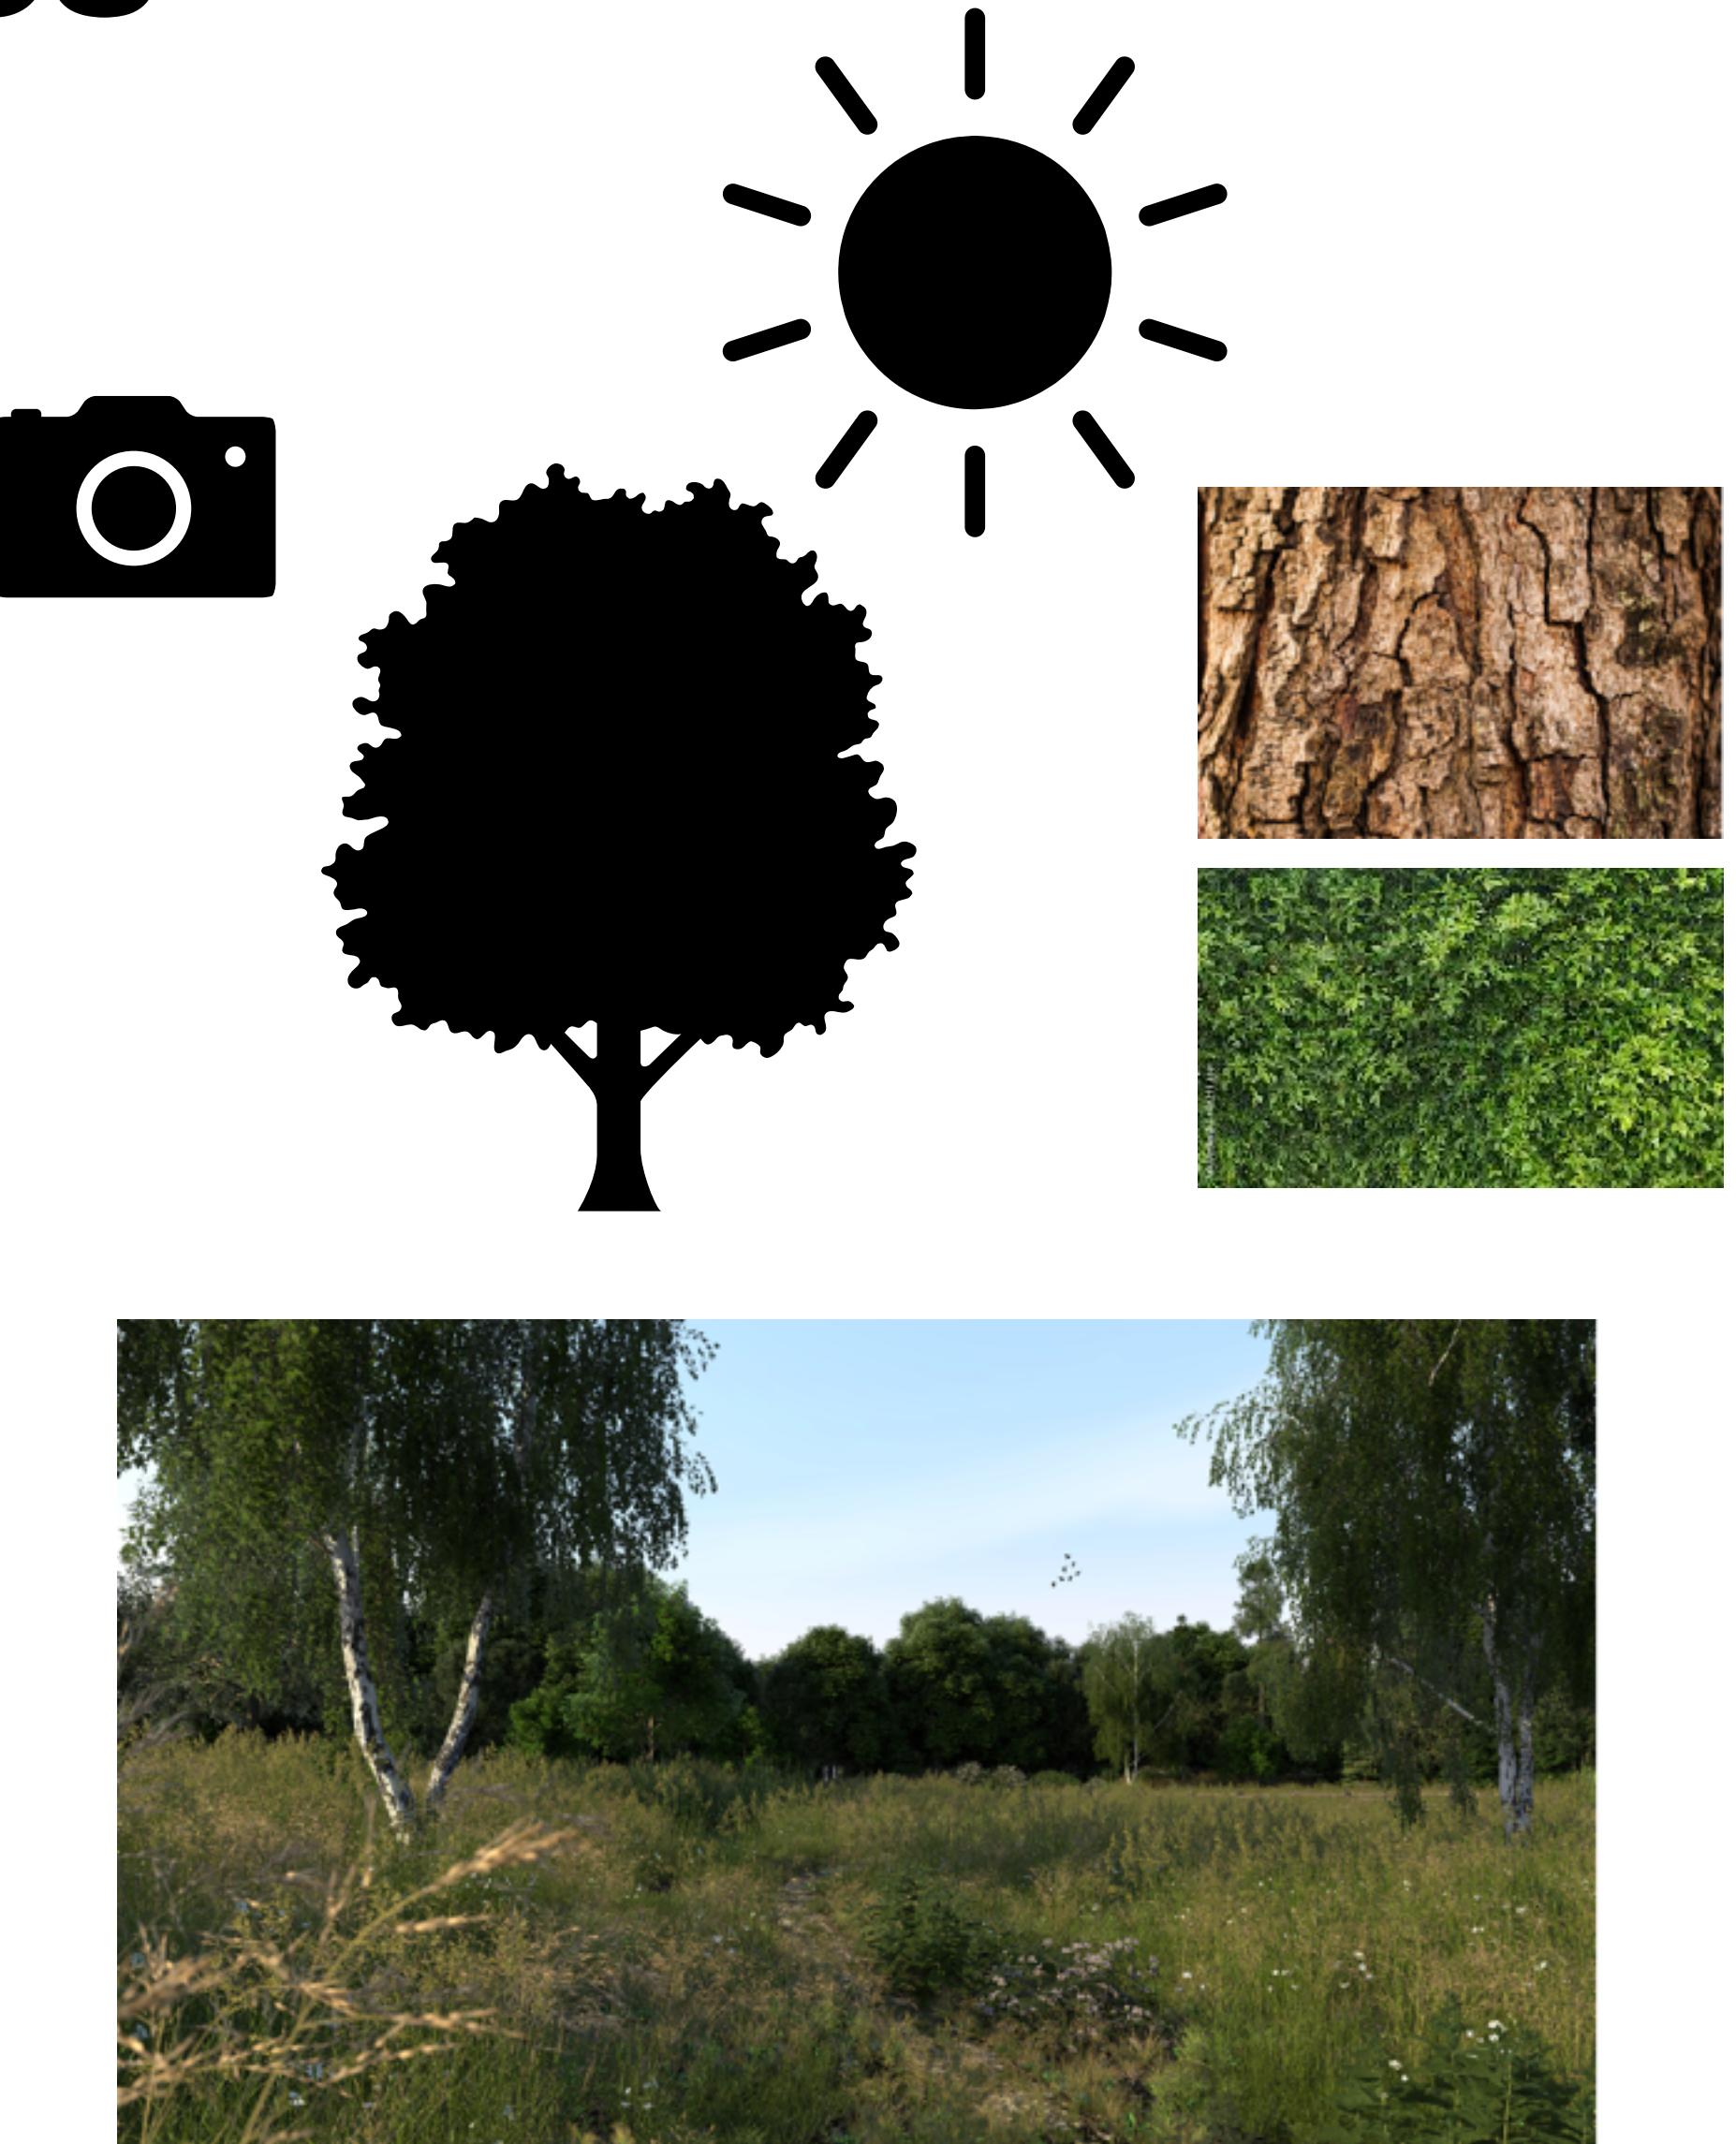
https://en.wikipedia.org/wiki/Stanford_bunny

<http://l2program.co.uk/786/little-tweaks-here-and-there>

Meshes & 3D Scenes

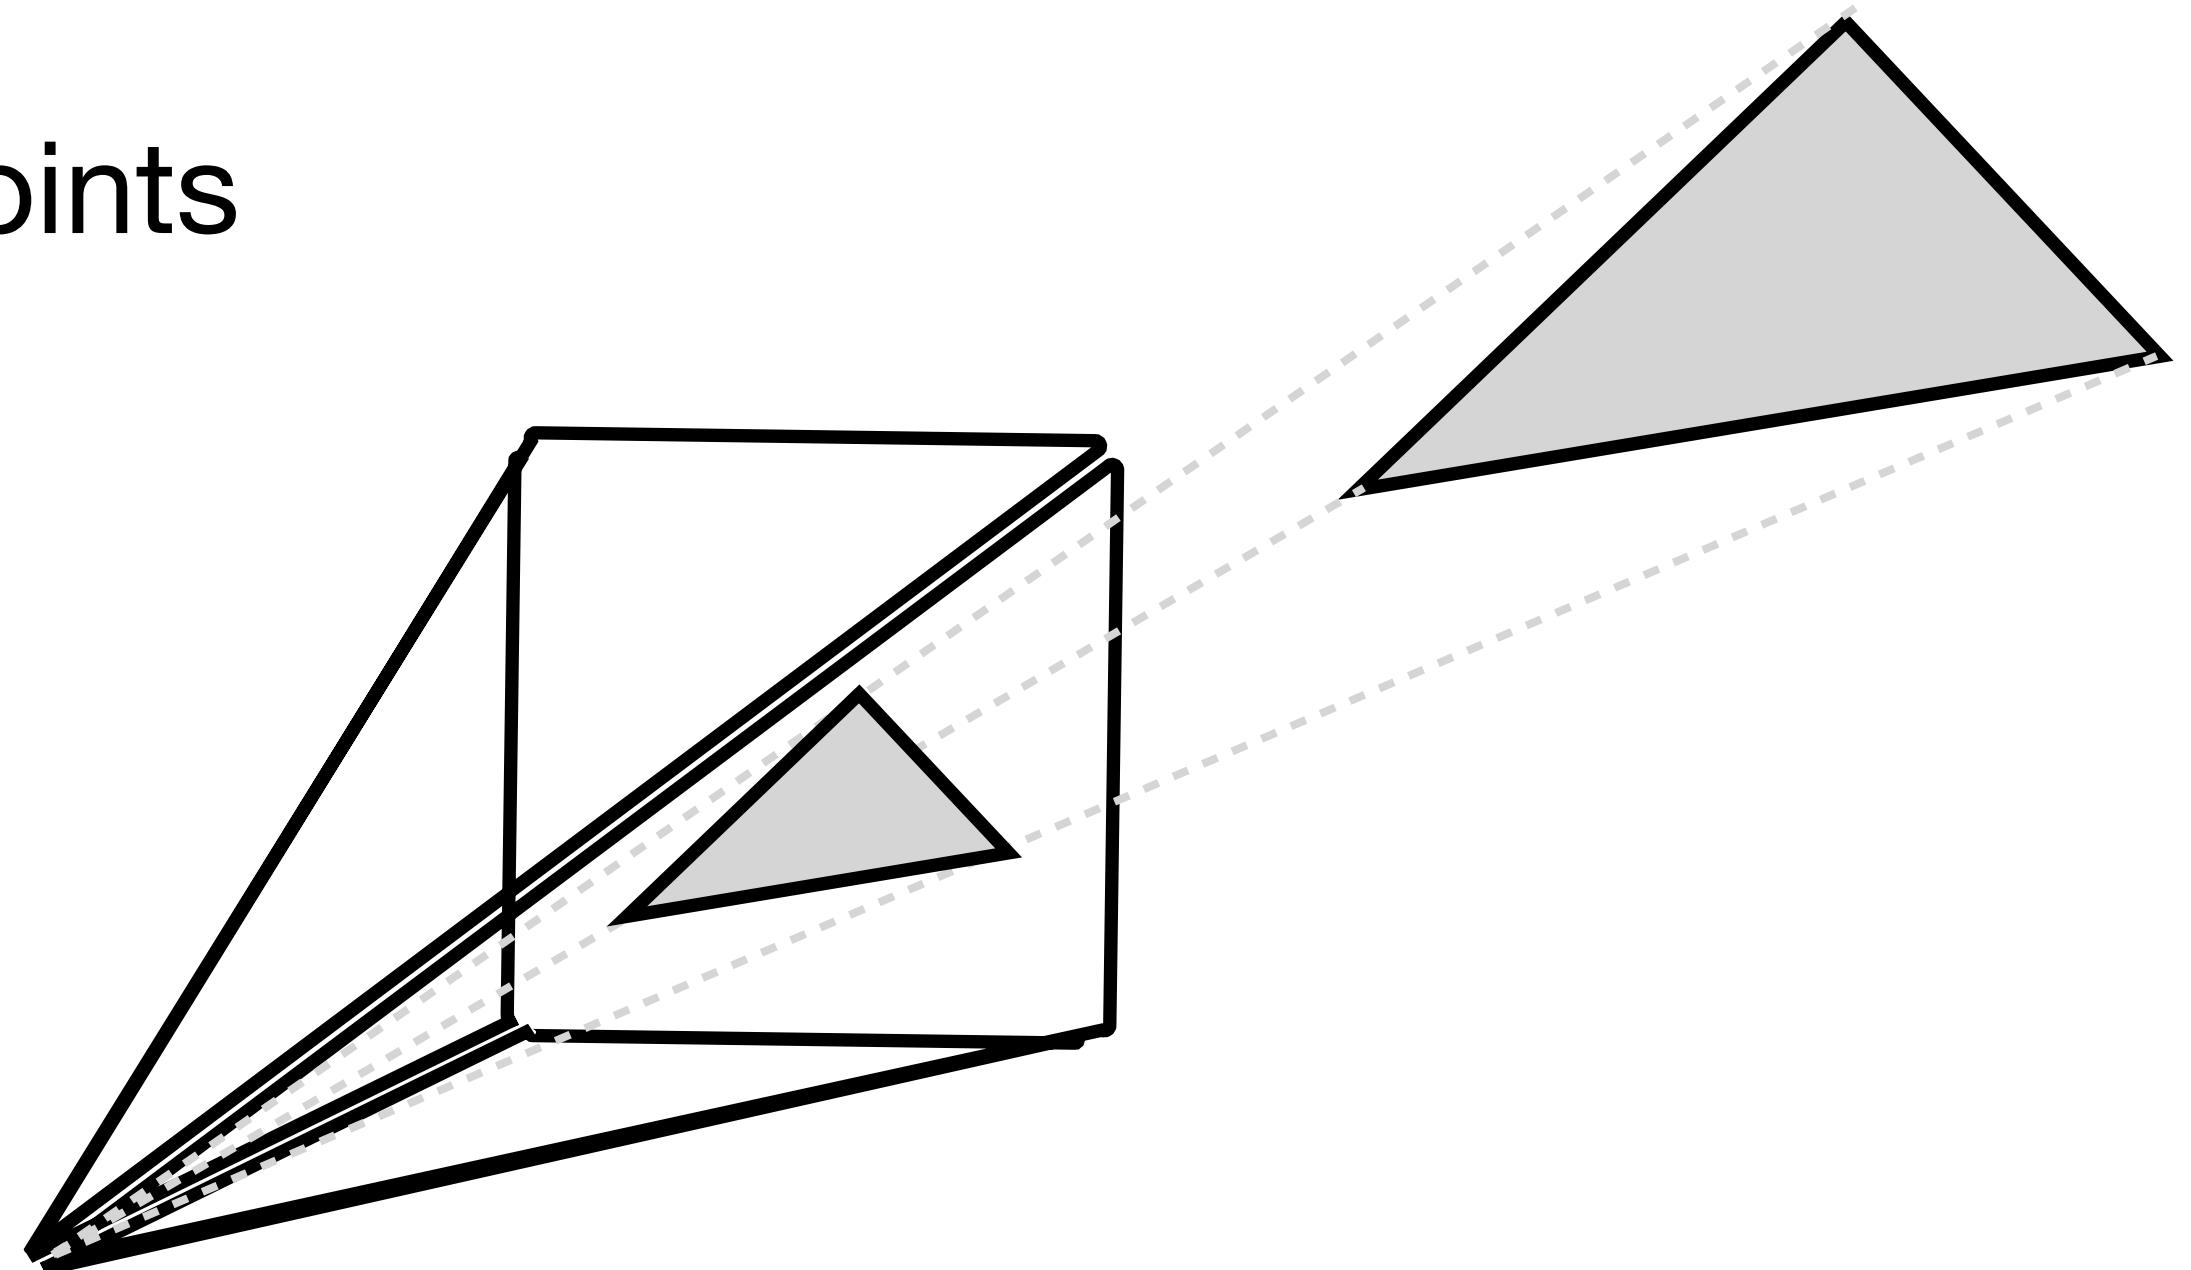
Modern Computer Graphics

- Scenes
 - Camera
 - Geometry
 - Textures
 - Light
- Renderer
 - Produces an image of a scene



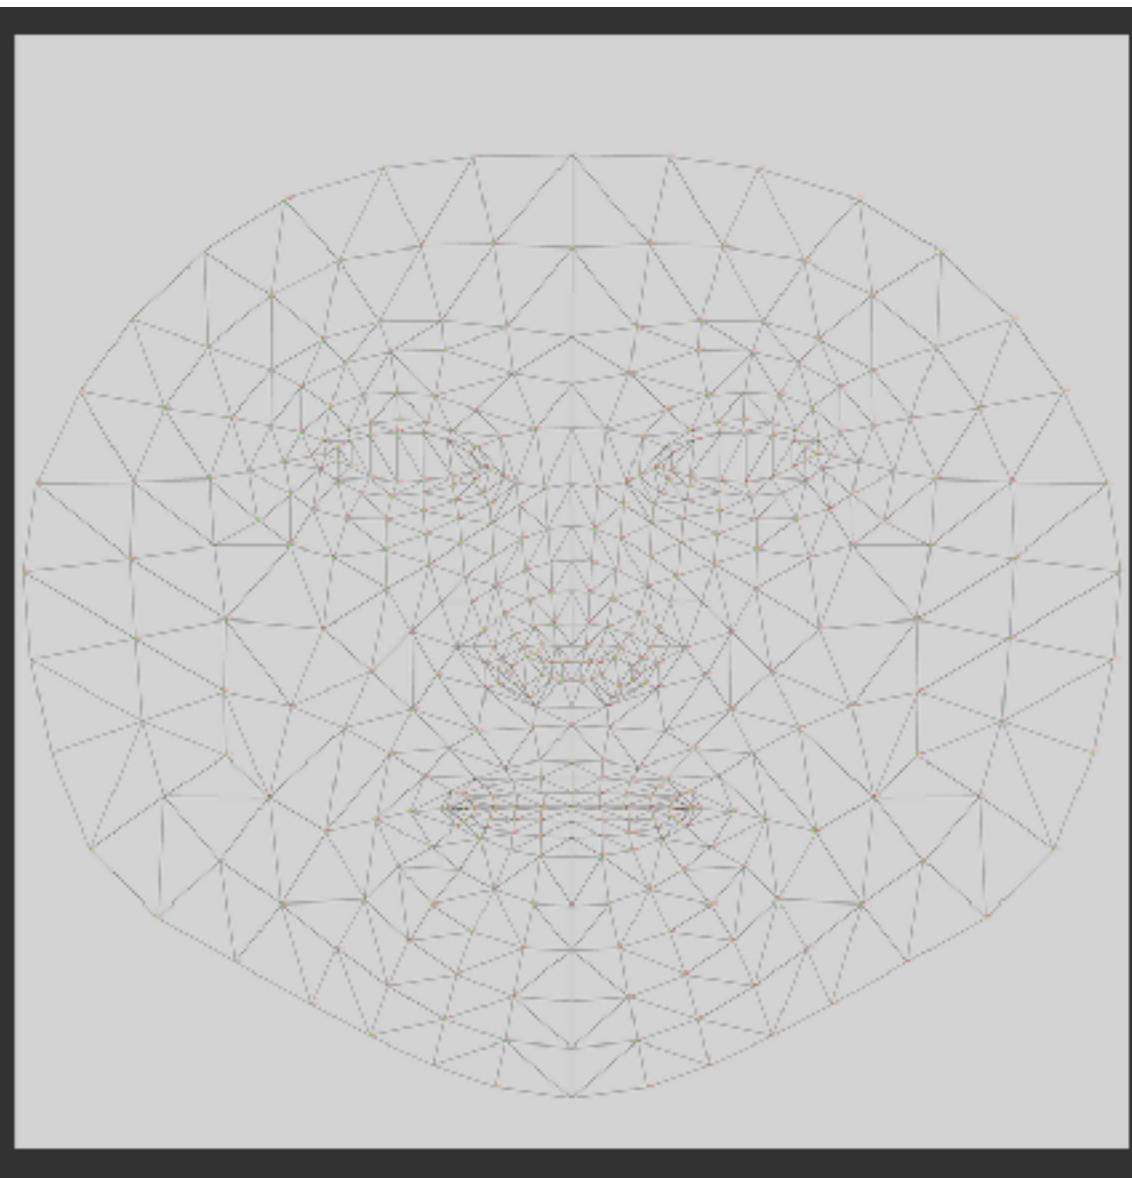
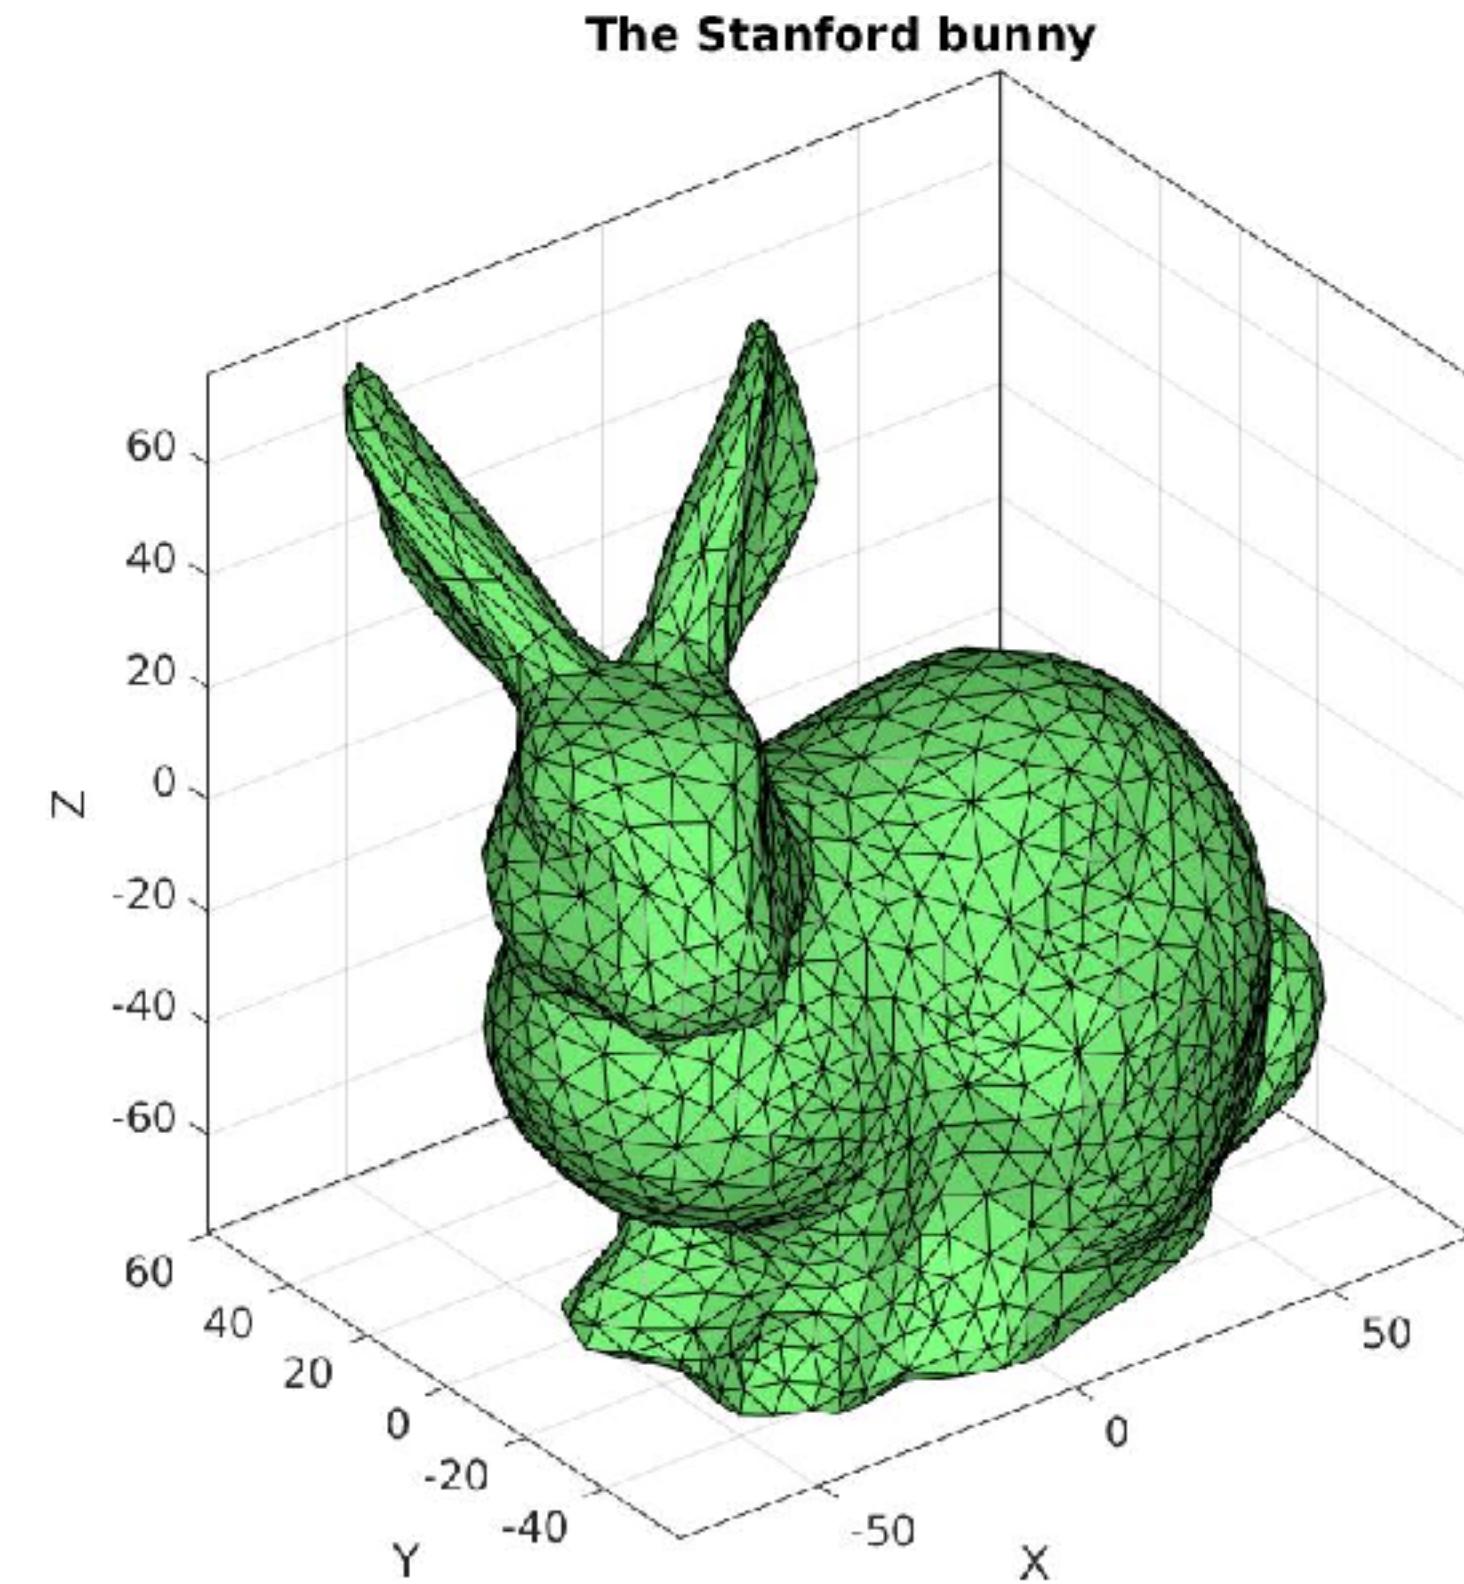
Triangles as a Primitive

- Multiple ways to represent shapes
 - Splines, level sets, polygonal meshes, points
- Triangles are very convenient
 - Three points always lie on a plane
 - Still a triangle after projection
 - Easy to tell if a point is inside

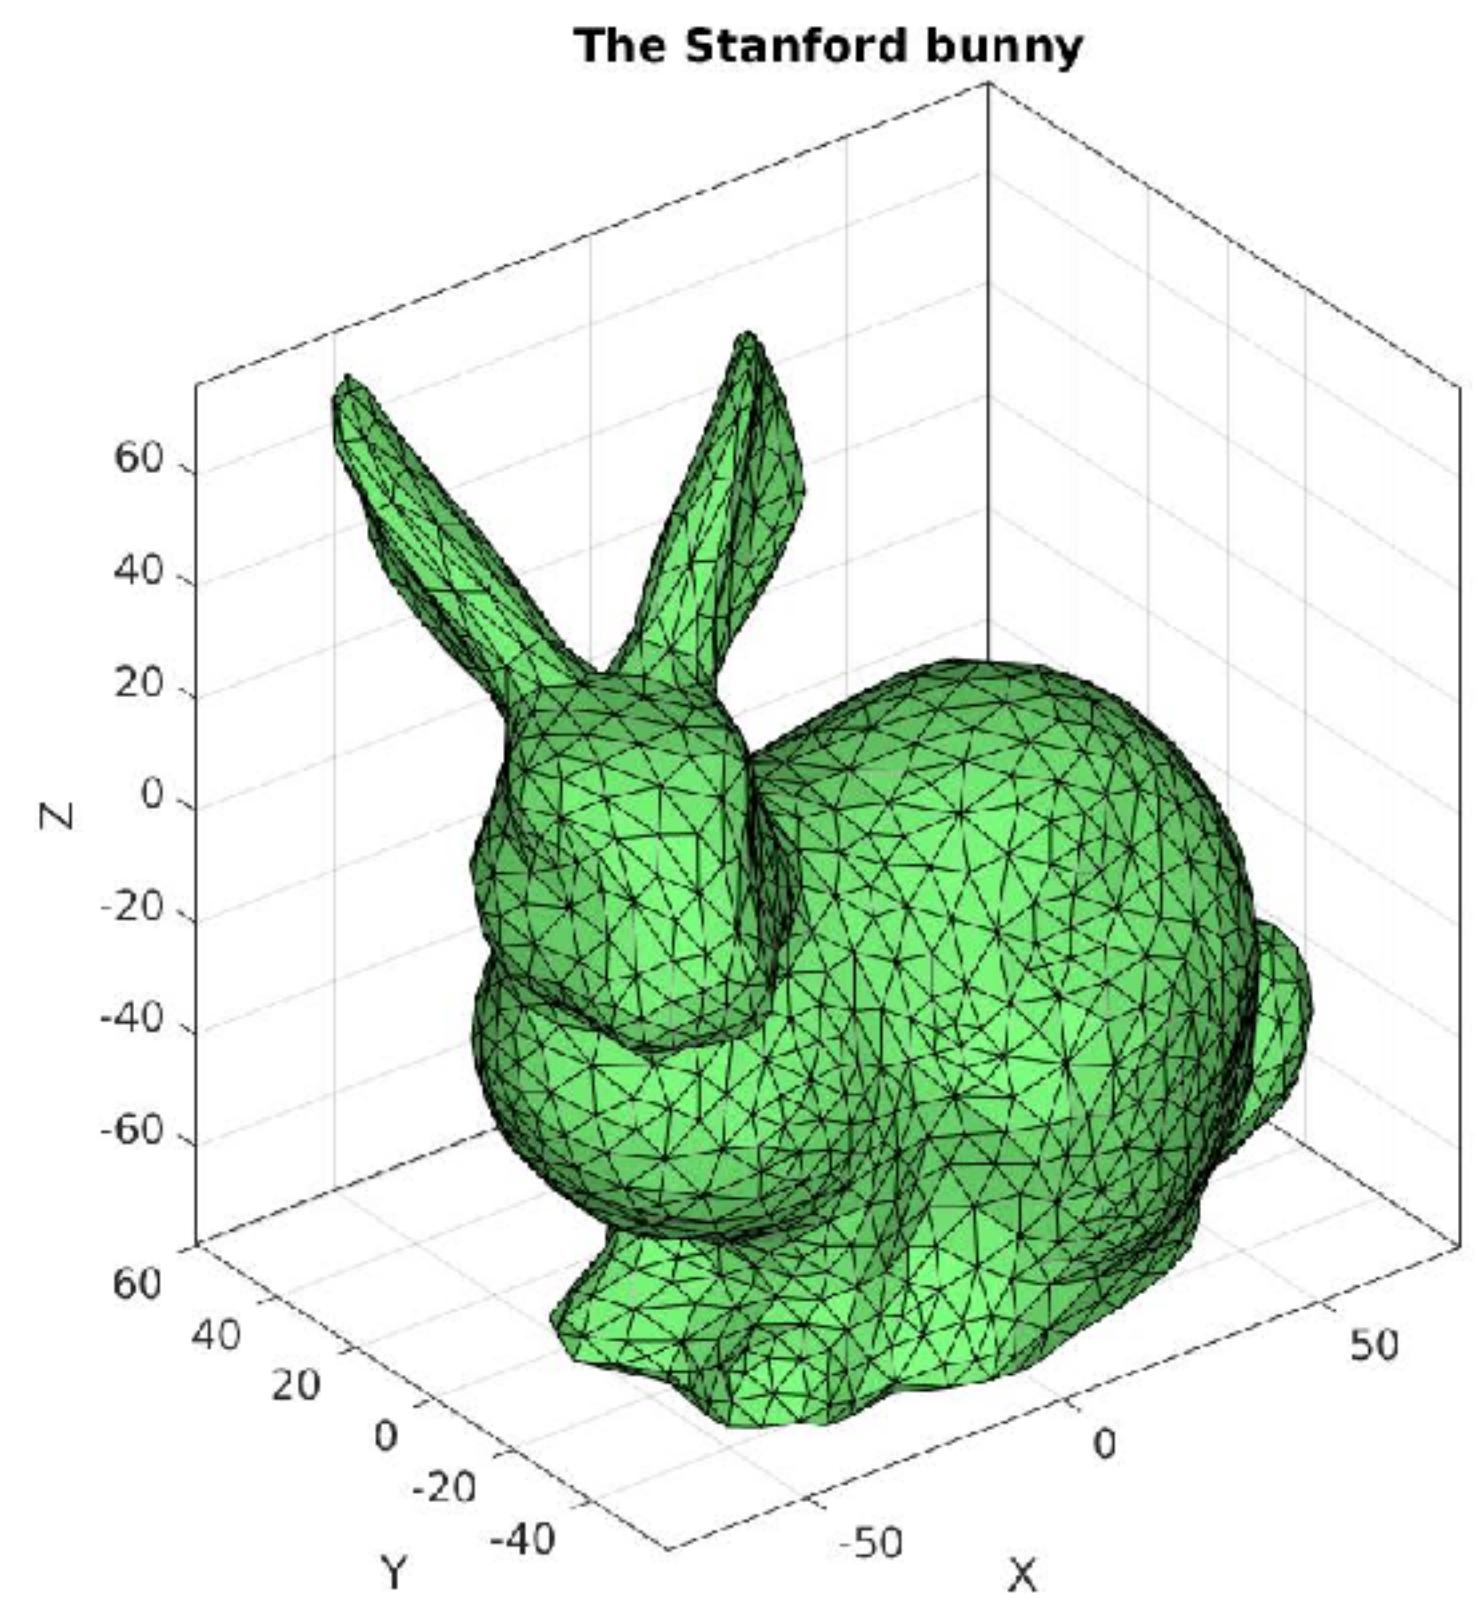


Polygonal Meshes Along with Textures

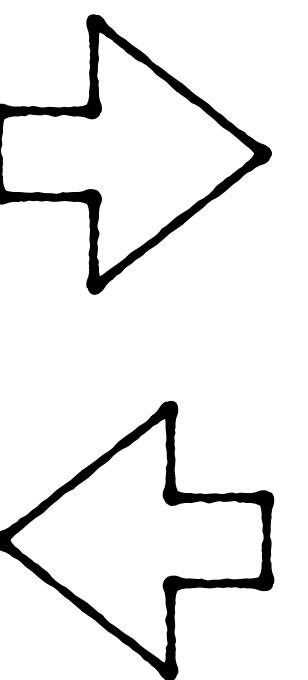
- Mesh consists of
 - Vertices $x_i \in \mathbb{R}^3$, $i = 1, \dots, N$
 - Triangles (p_{j0}, p_{j1}, p_{j2}) , $p_{ij} \in \{1, \dots, N\}$
- Texture
 - Function $T : [0,1]^2 \rightarrow \mathbb{R}^C$
 - UV-map: $f : \{1, \dots, N\} \rightarrow [0,1]^2$



Physically Based Rendering



graphics

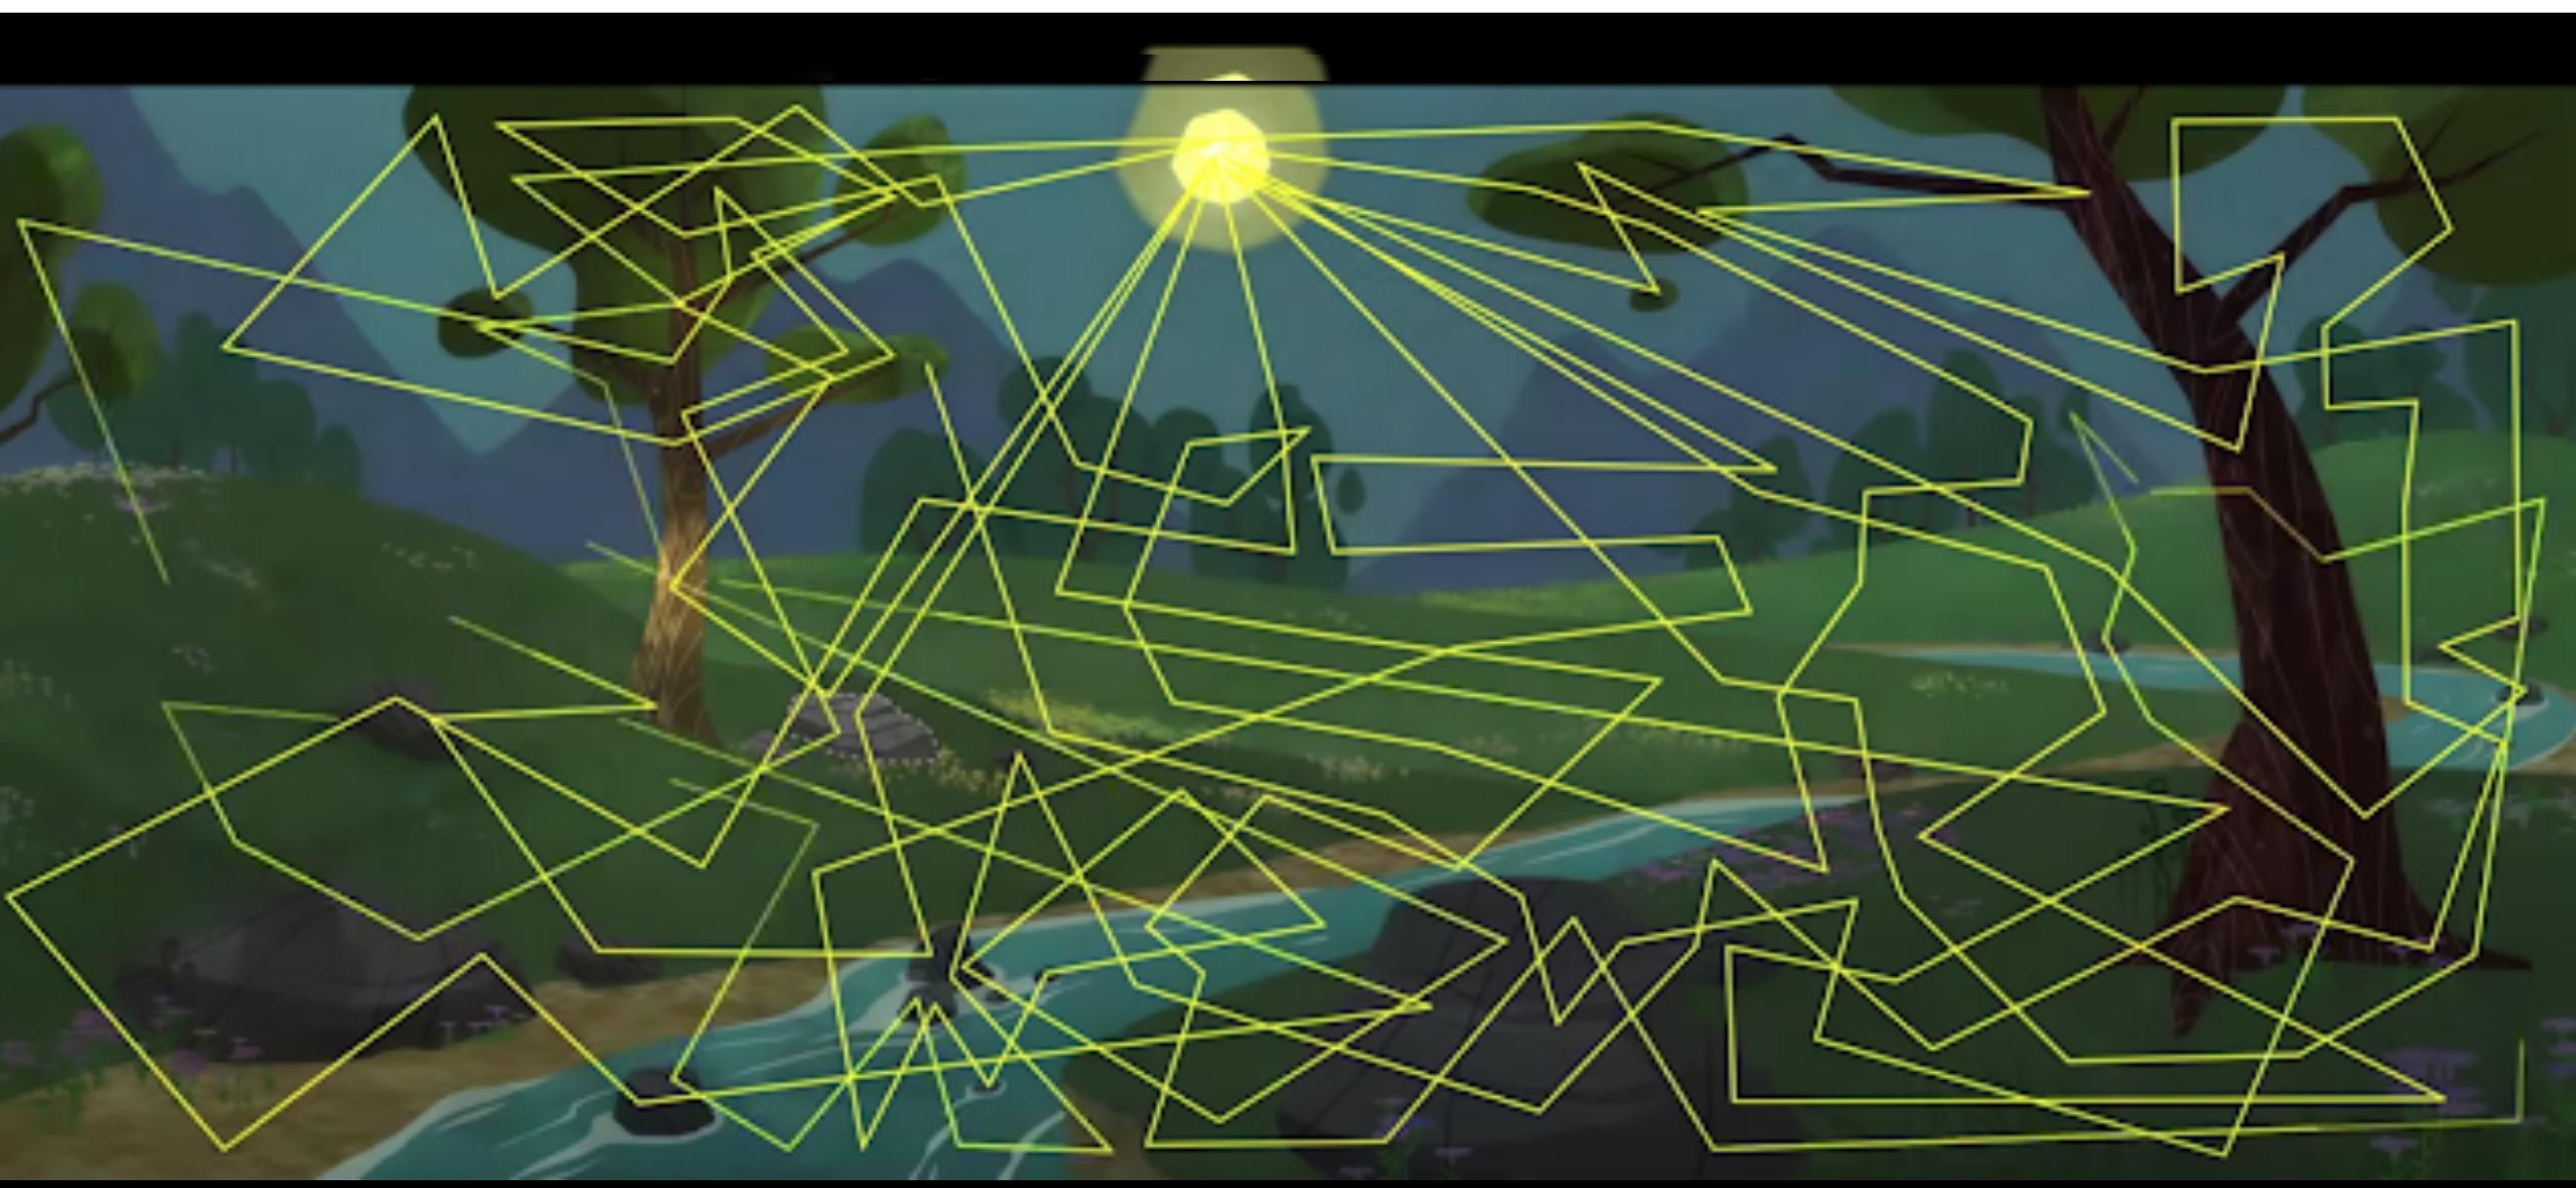


vision



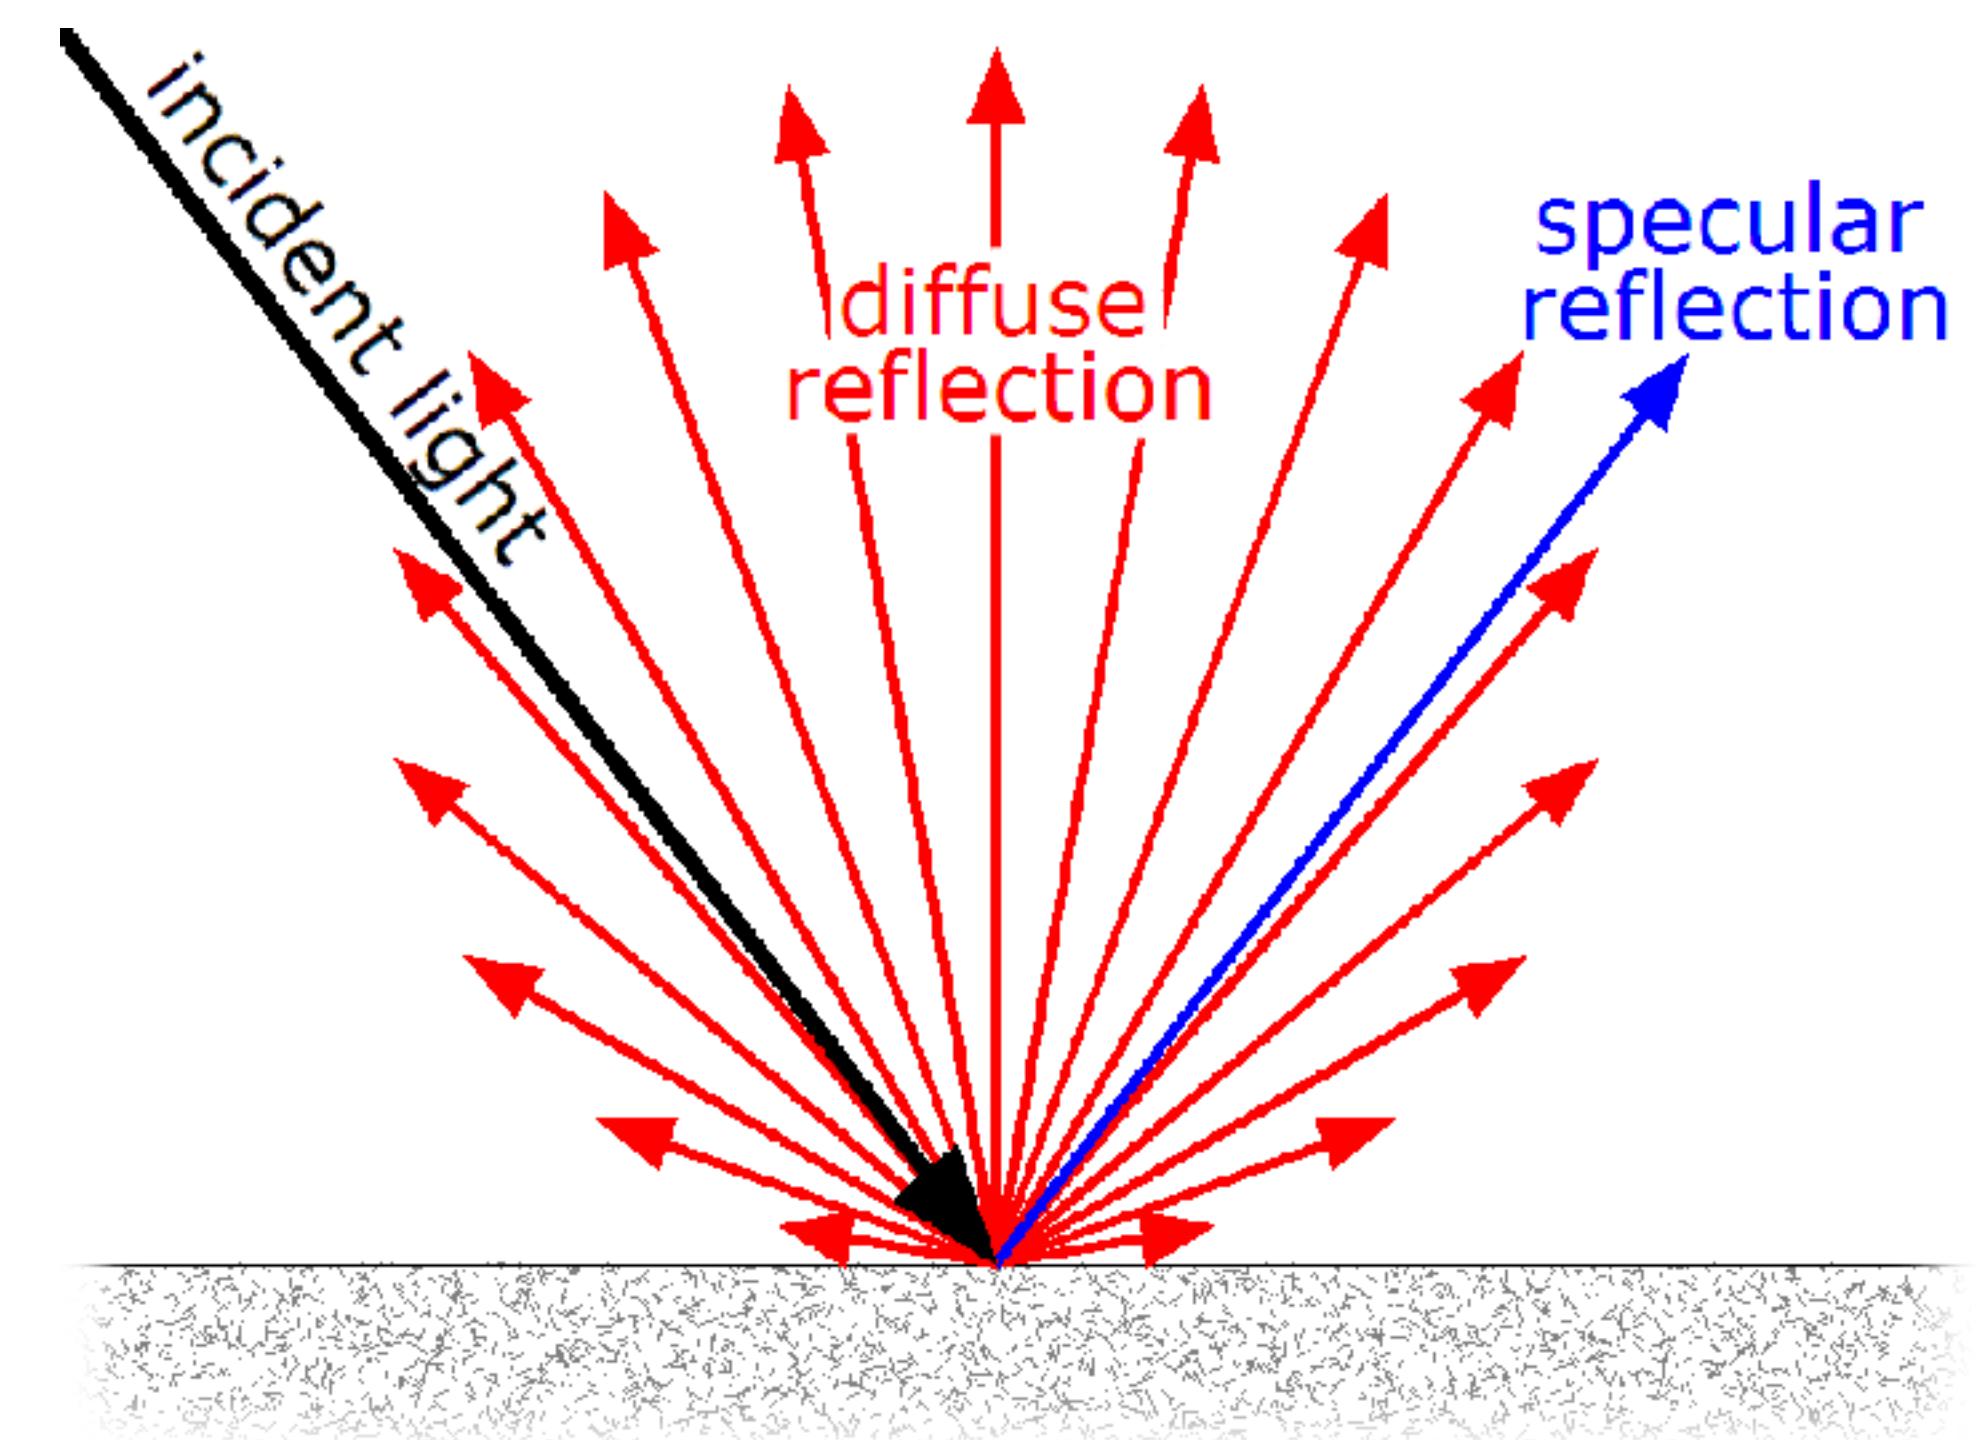
Path Tracing

- Idea: physical simulation of light propagation
- Emit *multiple* rays
- Rays interact with scene geometry
- Register rays that reach camera sensor



Light Interaction

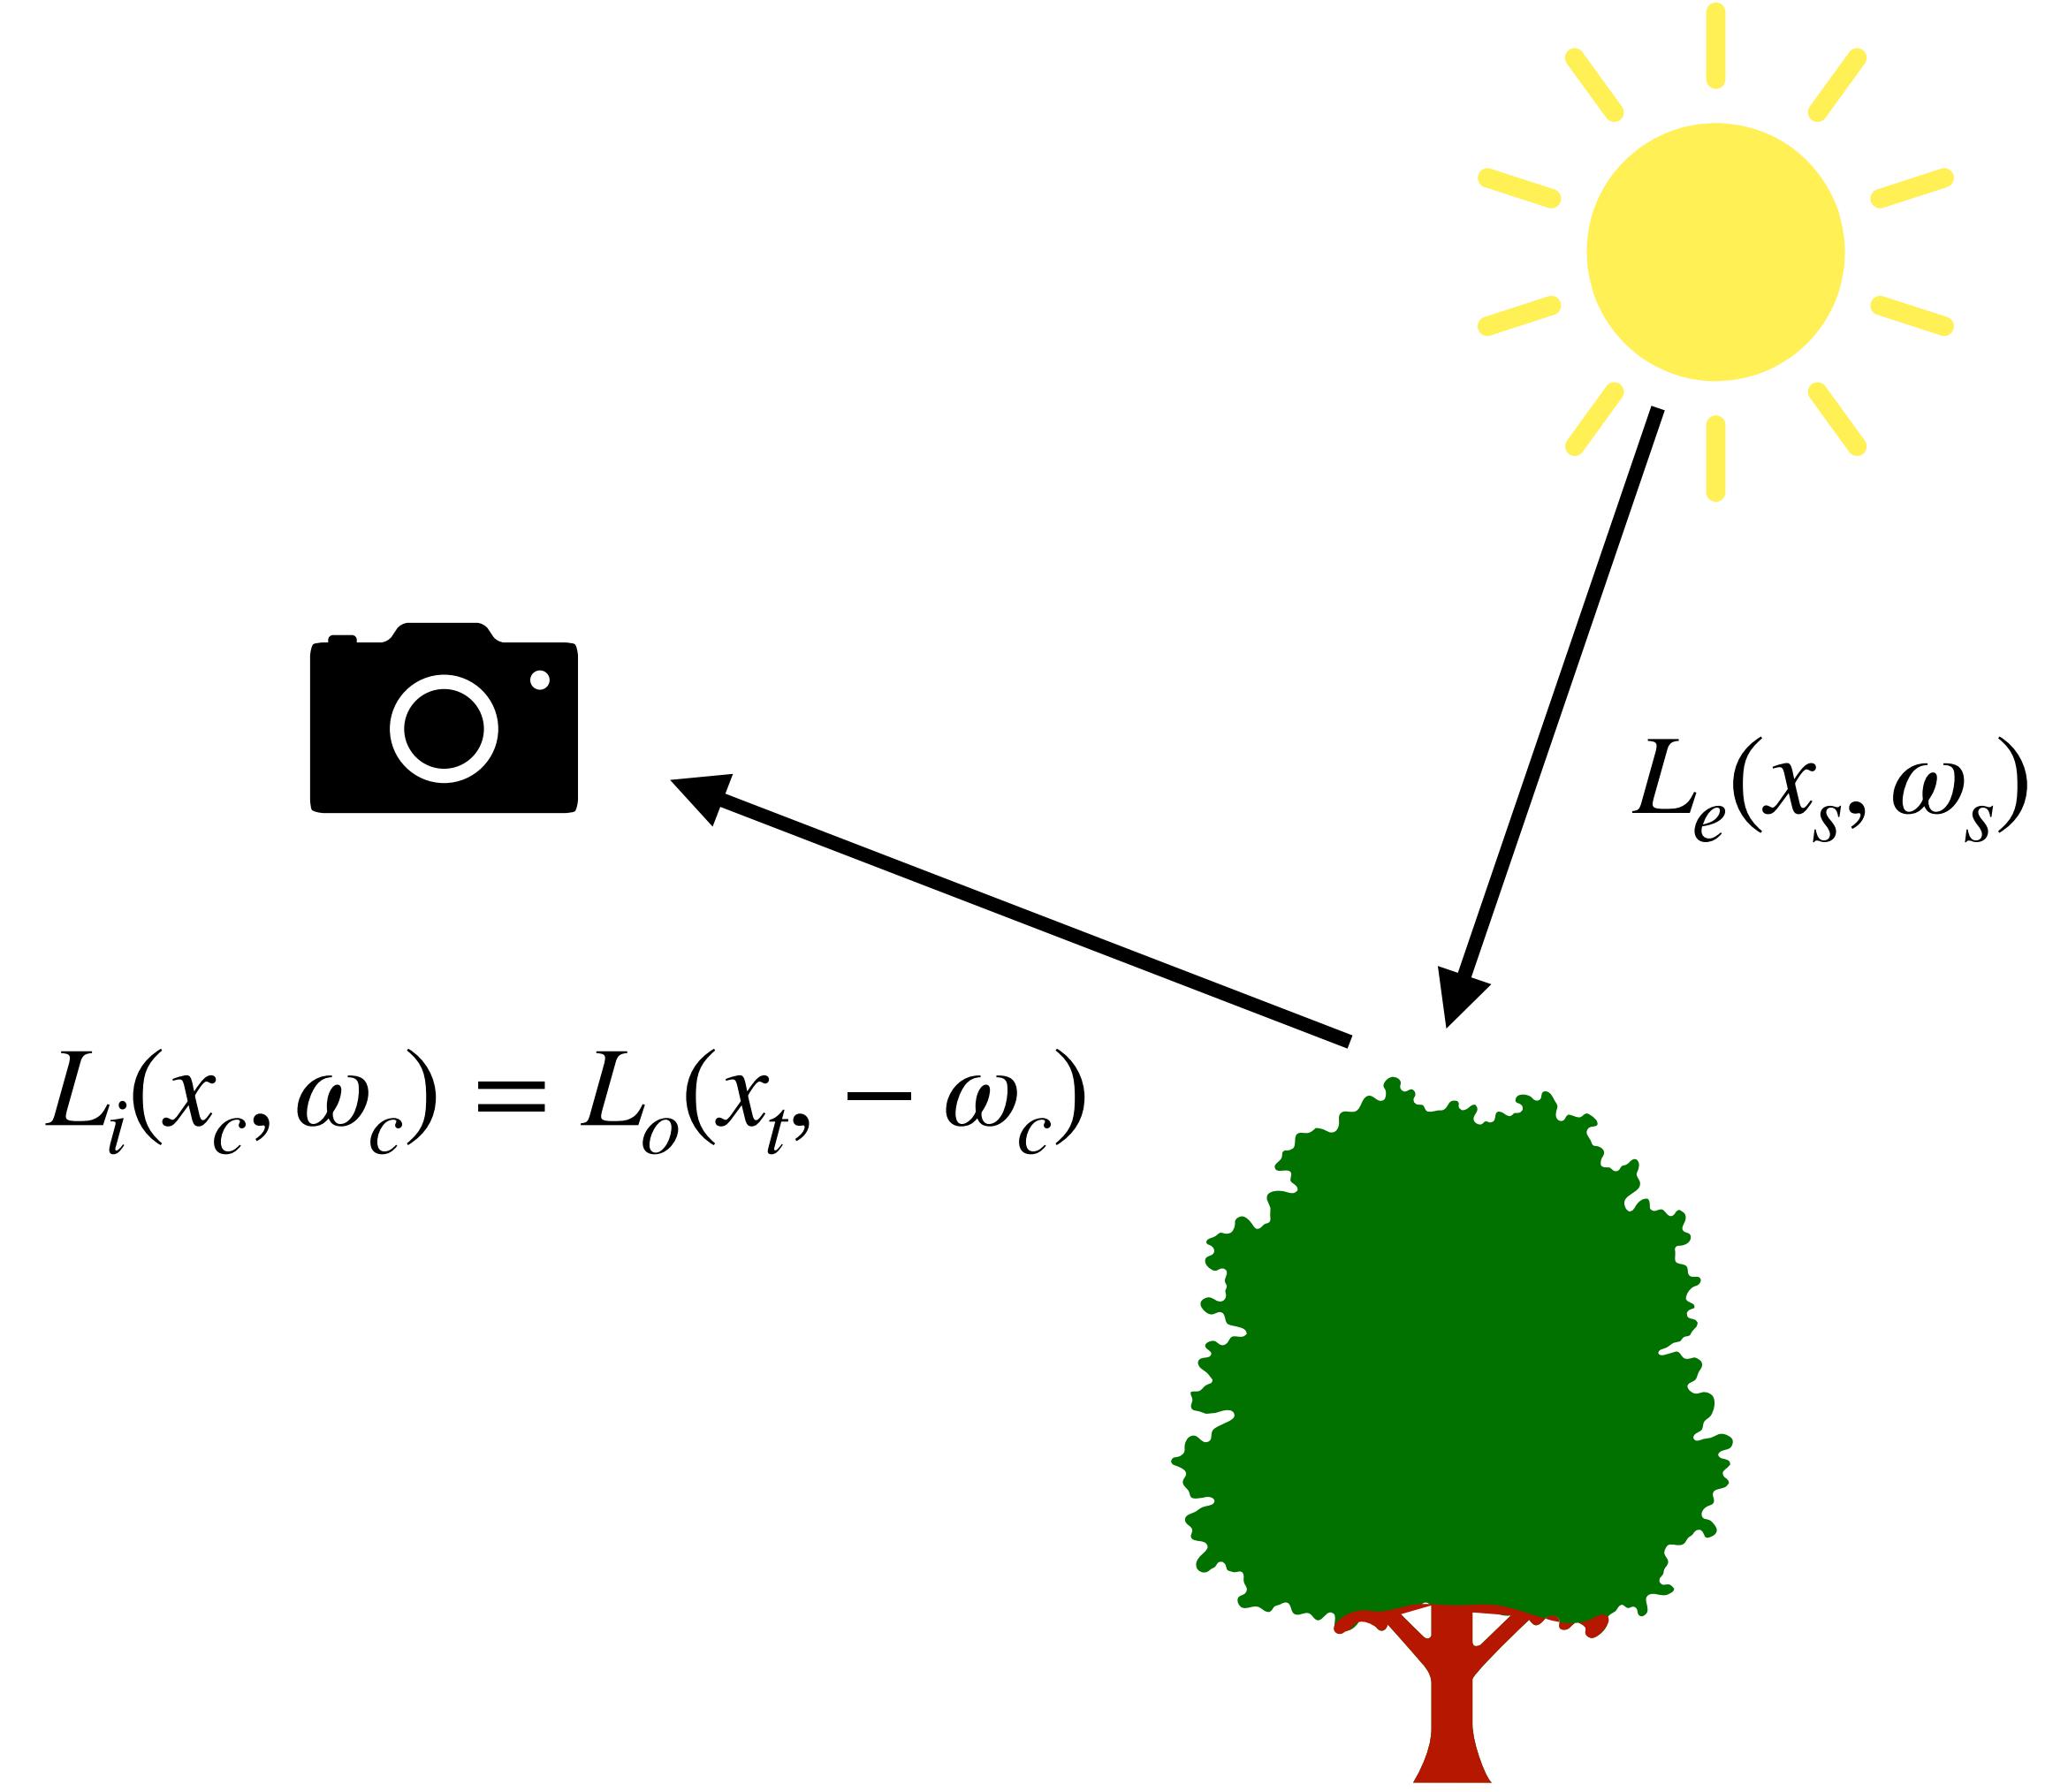
- Diffuse reflection
 - All reflection directions are equal
- Specular reflection
 - Perfect mirror
- In general: BRDF
 - Bidirectional Reflectance **Distribution** Function
 - Function $f_r(\omega_i, \omega_r) : [0,1]^2 \times [0,1]^2 \rightarrow \mathbb{R}$



Radiance Functions

Aggregating Simulation Results

- Three kinds of radiance
 - Emitted $L_e(x, \omega_o)$
 - Outgoing $L_o(x, \omega_o)$
 - Incoming $L_i(x, \omega_i)$
- Rendering amounts to computing $L_i(x, \omega)$
 - Point x being the camera location
 - Angle grid $\omega \in \Omega$ covering the virtual image plane



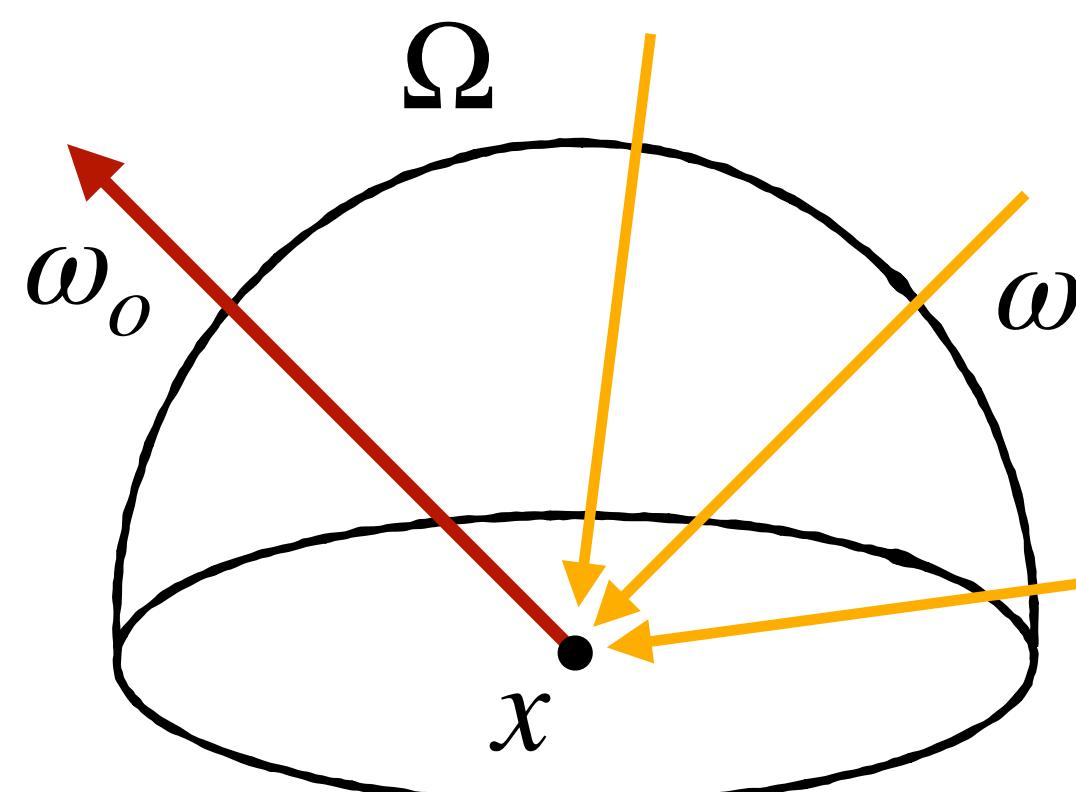
Rendering Equation

Framework Behind Path Tracing

$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{\Omega} f(x, \omega_i, \omega_o) L_i(x, \omega_i) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Diagram illustrating the rendering equation framework:

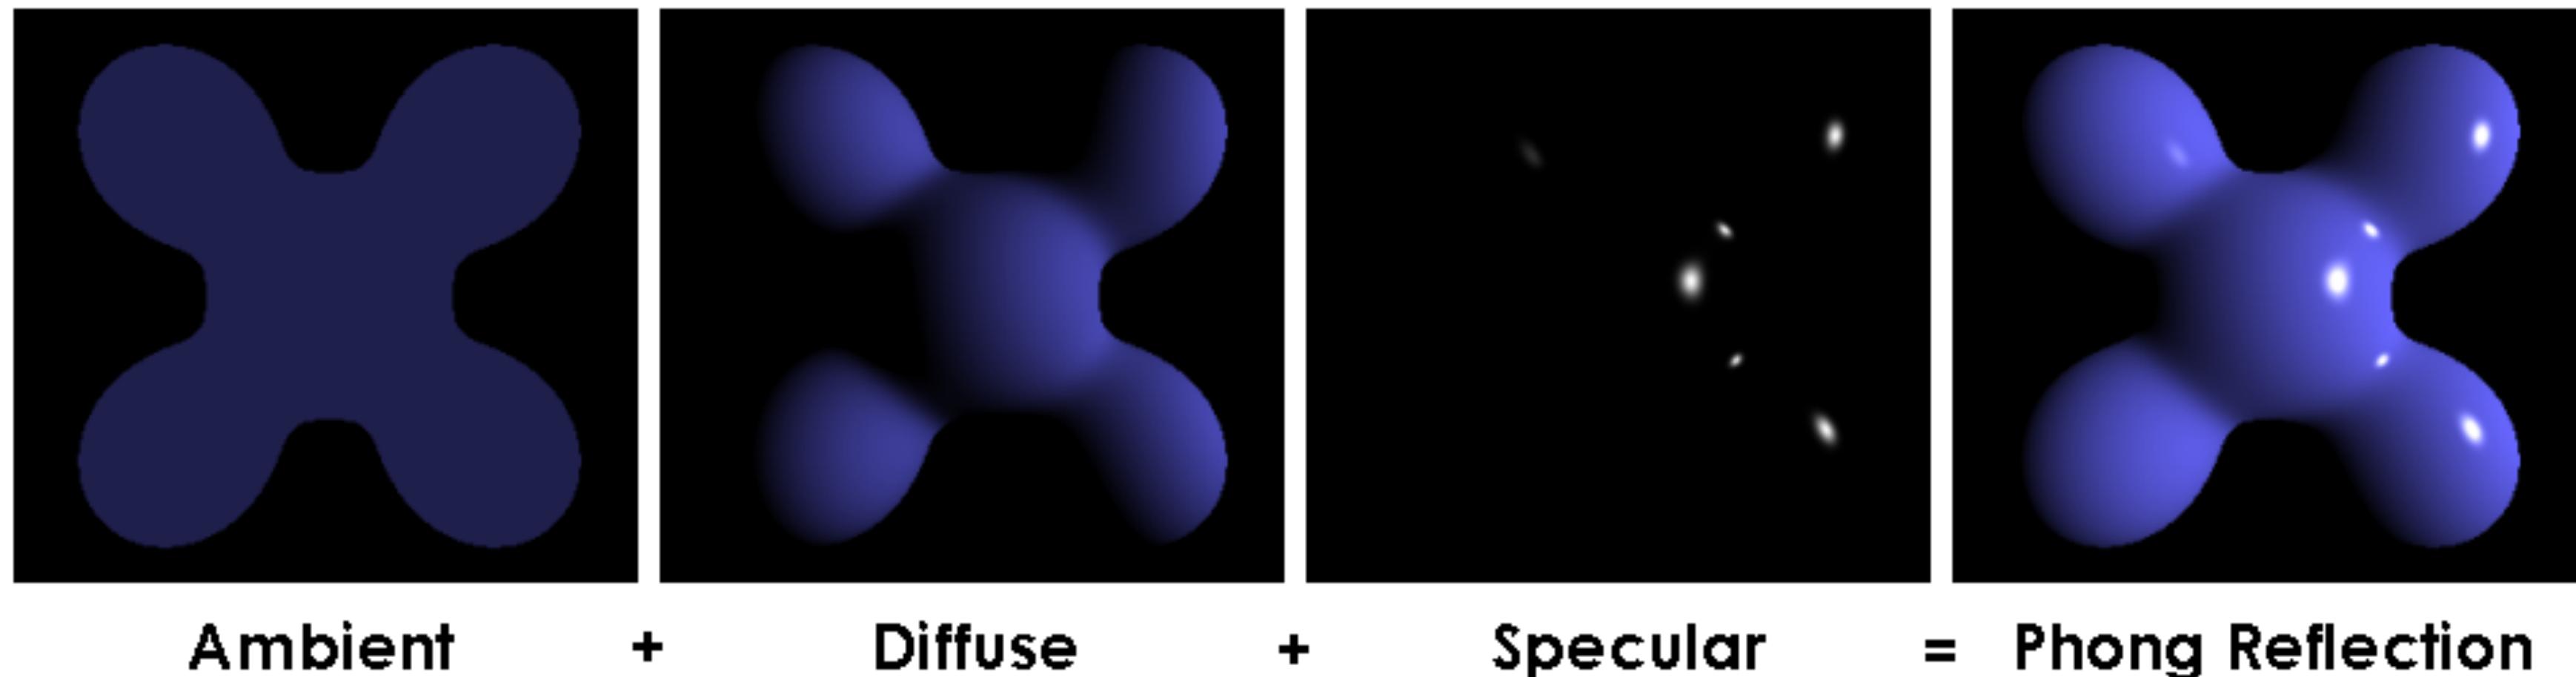
- emitted radiance in ω_o** : A dashed arrow points down to the term $L_e(x, \omega_o)$.
- outgoing radiance in ω_o** : A dashed arrow points up to the term $L_o(x, \omega_o)$.
- incoming radiance**: A dashed arrow points down to the integral term.
- angle weight of photon density**: A dashed arrow points up to the integral term.
- fraction of light that is reflected into ω_o** : A dashed arrow points up to the term $f(x, \omega_i, \omega_o)$.



Side Note on PBR Textures

Two Kinds of Reflections

- Ambient texture $T_a \in \mathbb{R}^{H \times W \times 3}$
- Diffuse texture $T_d \in \mathbb{R}^{H \times W \times 3}$
- Specular Texture (metallic & roughness) $T_s \in \mathbb{R}^{H \times W \times 2}$



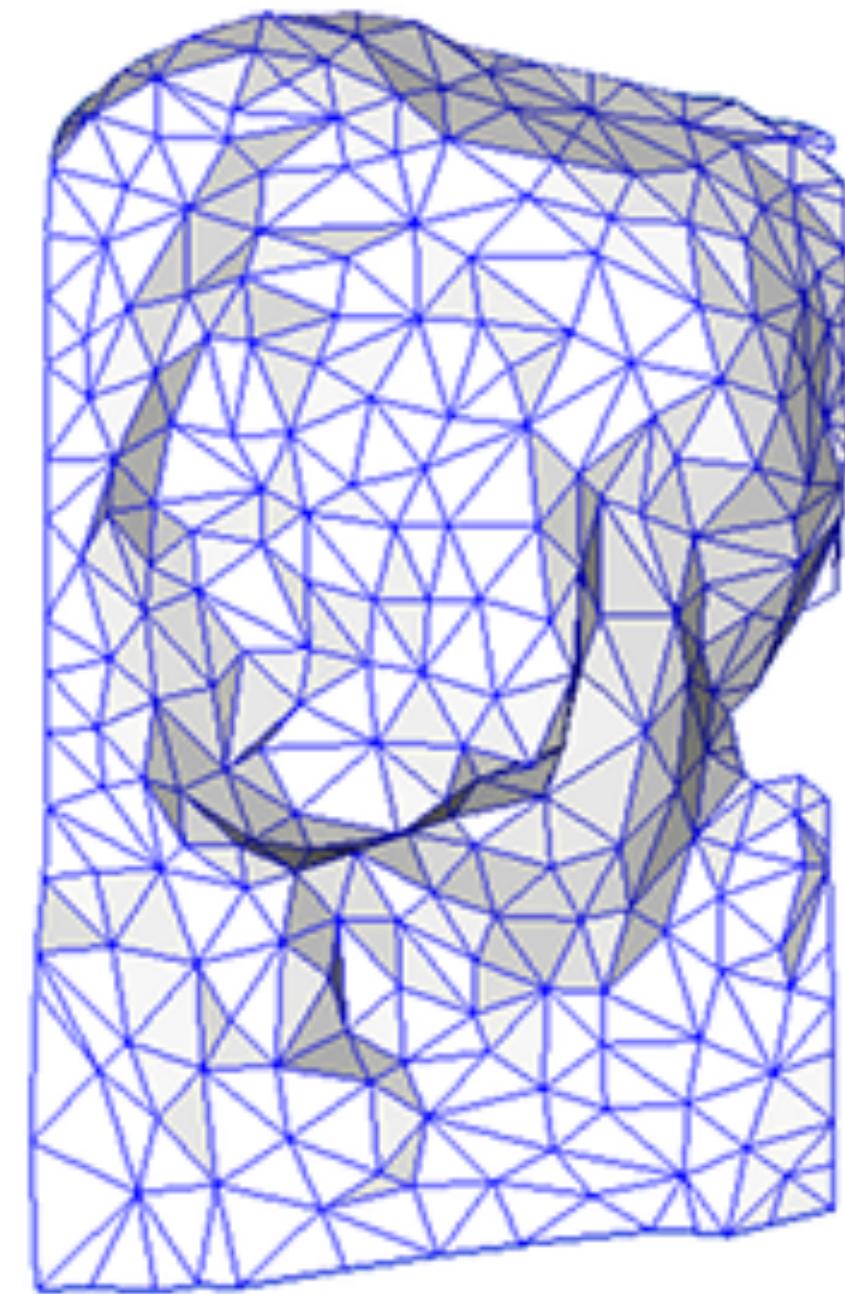
Side Note on PBR Textures

Normal Maps

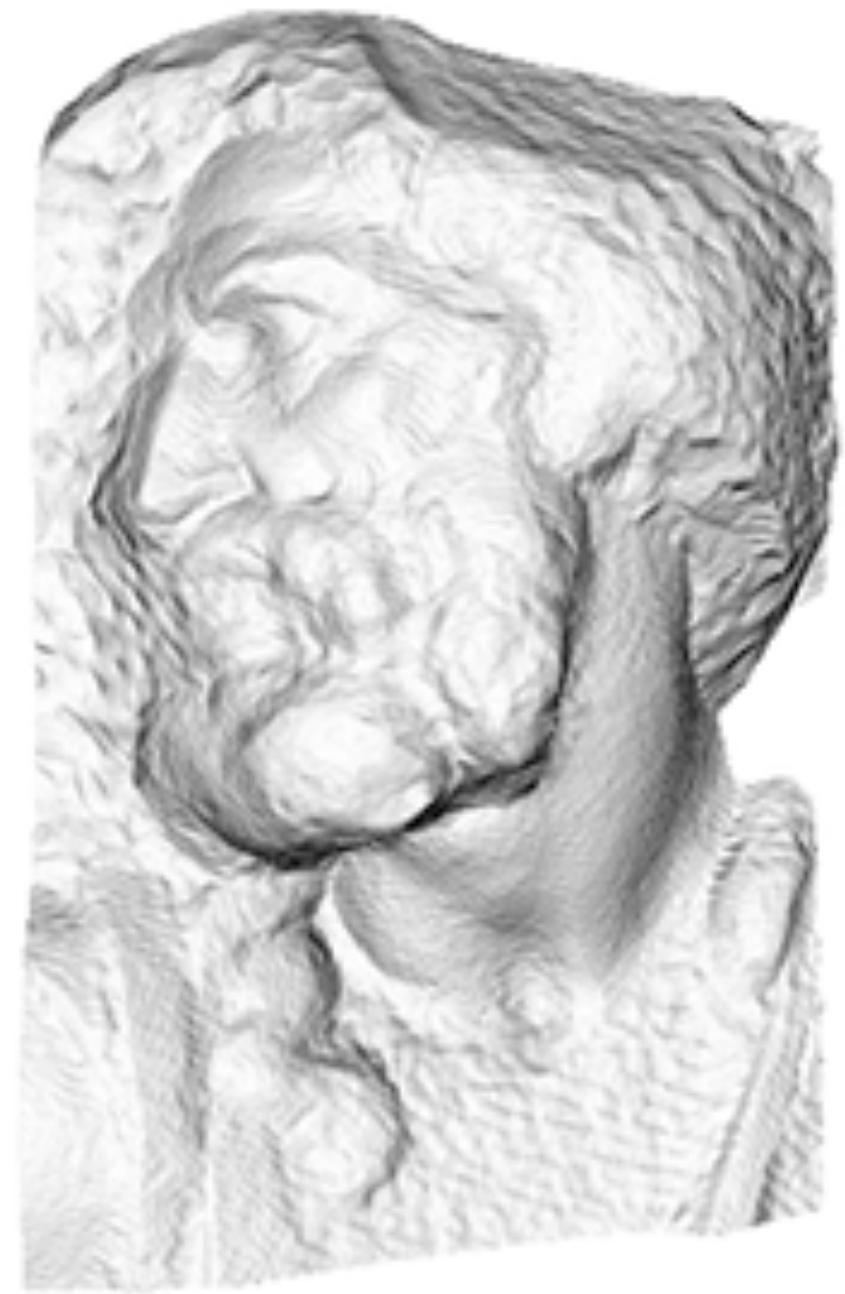
- Normal map $T_n \in \mathbb{R}^{H \times W \times 2}$



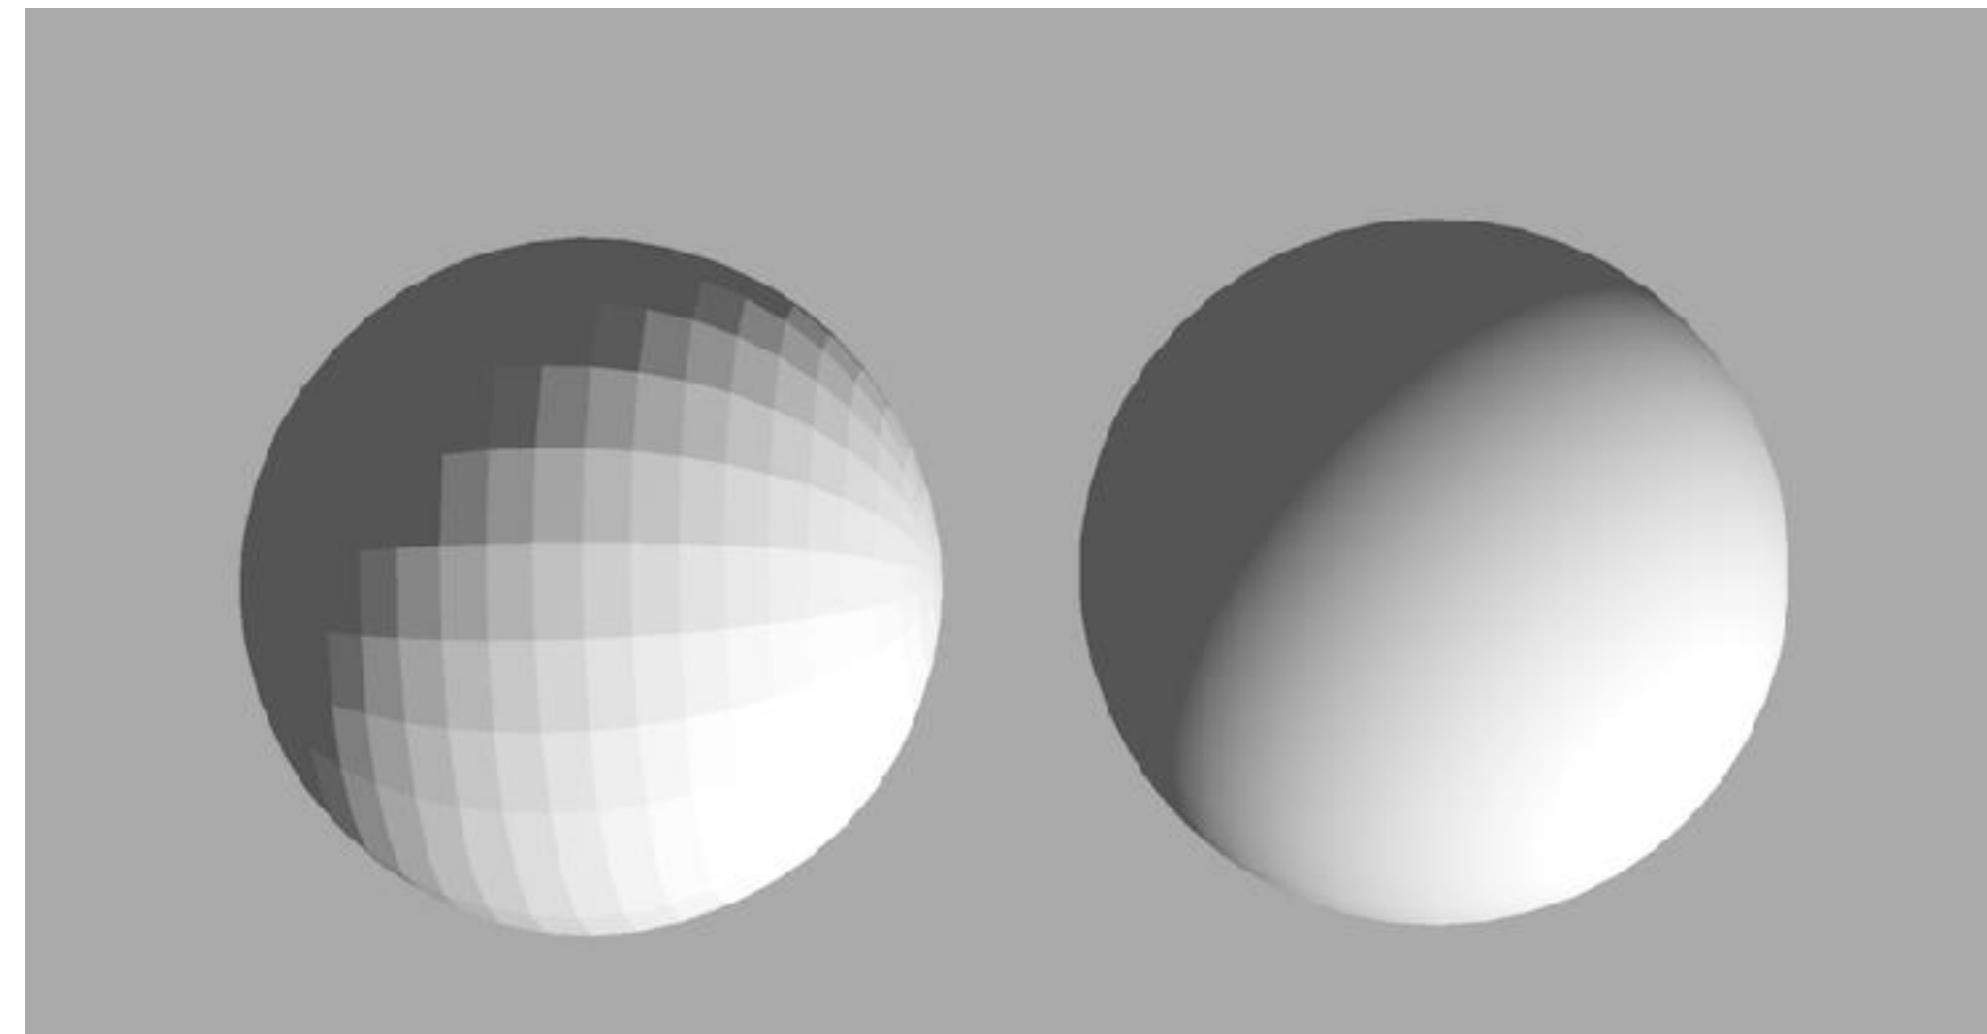
original mesh
4M triangles



simplified mesh
500 triangles



simplified mesh
and normal mapping
500 triangles



Solving Rendering Equation

Naive Rendering Pseudo-Code

for each pixel (x, y)

 compute **ray** origin r_o and direction r_d

for each object O in scene

if ray intersect O and is closest for far

 save O_{closest}

 compute emitted radiance at intersection with O_{closest}

 recursively compute incoming radiance at the intersection

Implementations

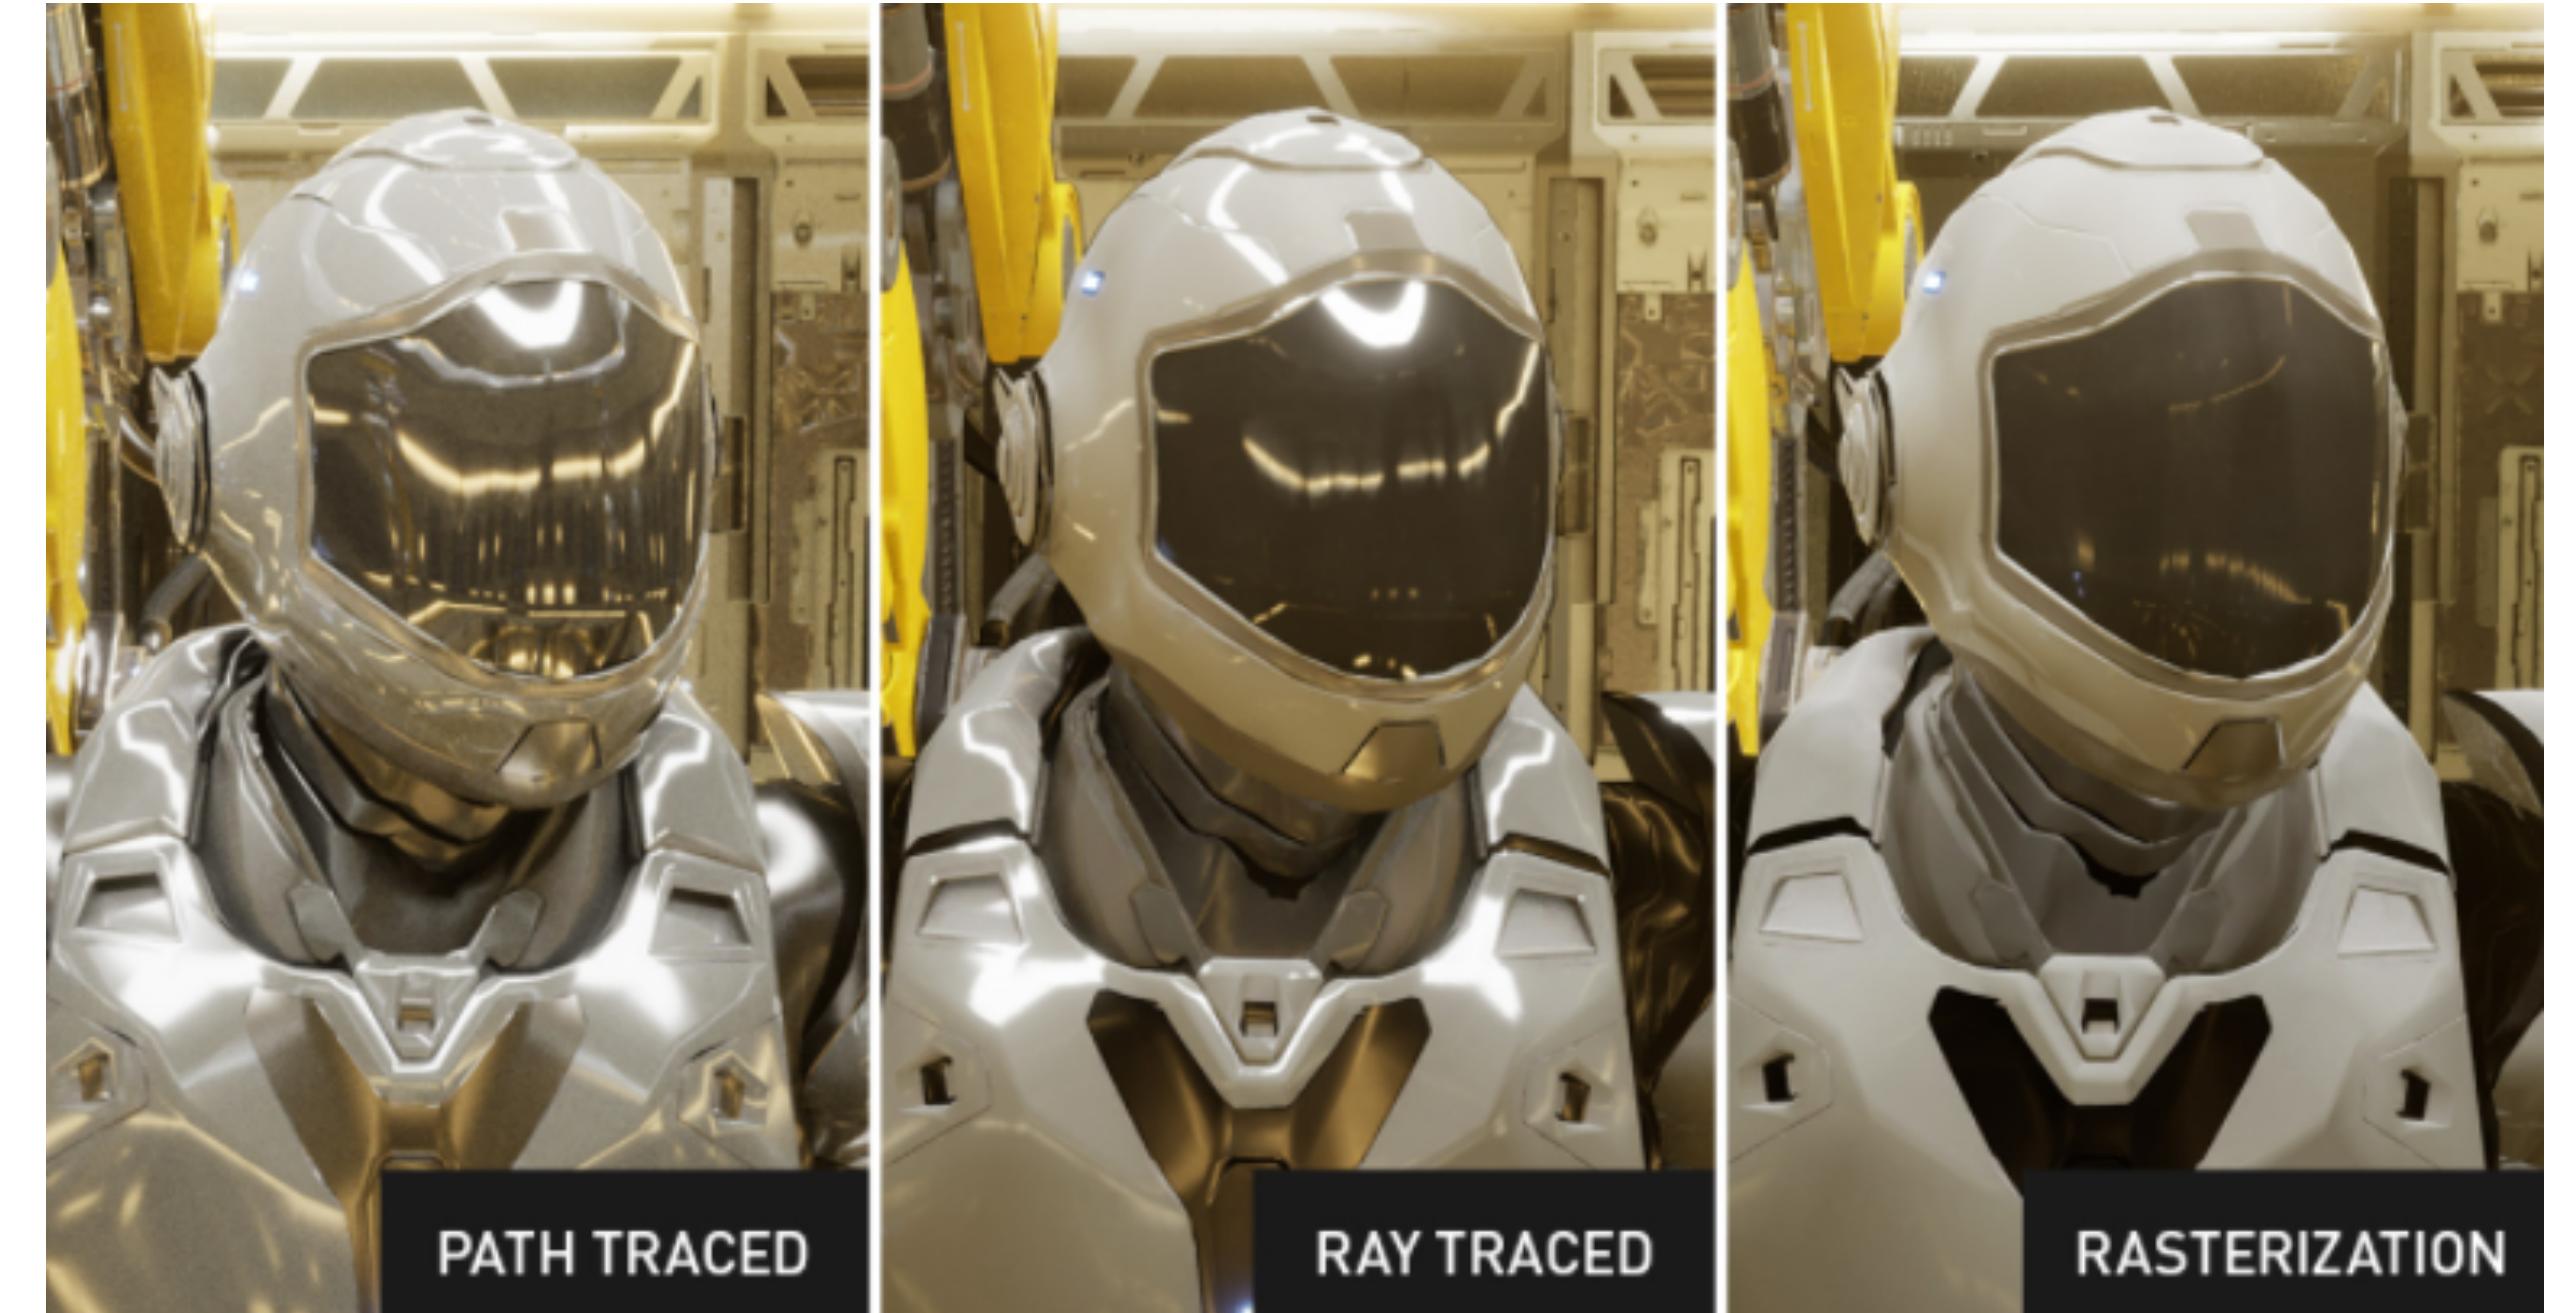
- Simple implementation in JAX & more detailed explanation
 - <https://blog.evjang.com/2019/11/jaxpt.html>
- Mitsuba Renderer
 - <https://www.mitsuba-renderer.org/>
- Physically-based rendering in 3D reconstruction
 - <https://nvlabs.github.io/nvdiffrcmc/>



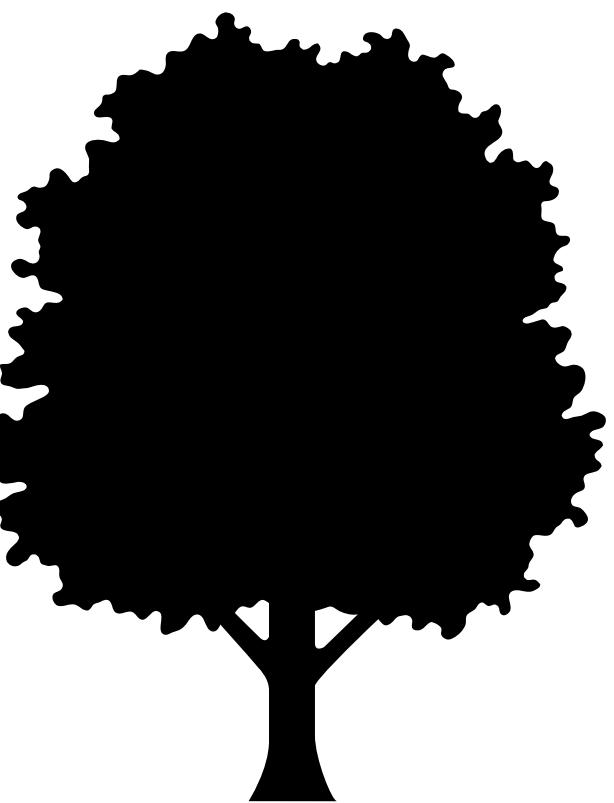
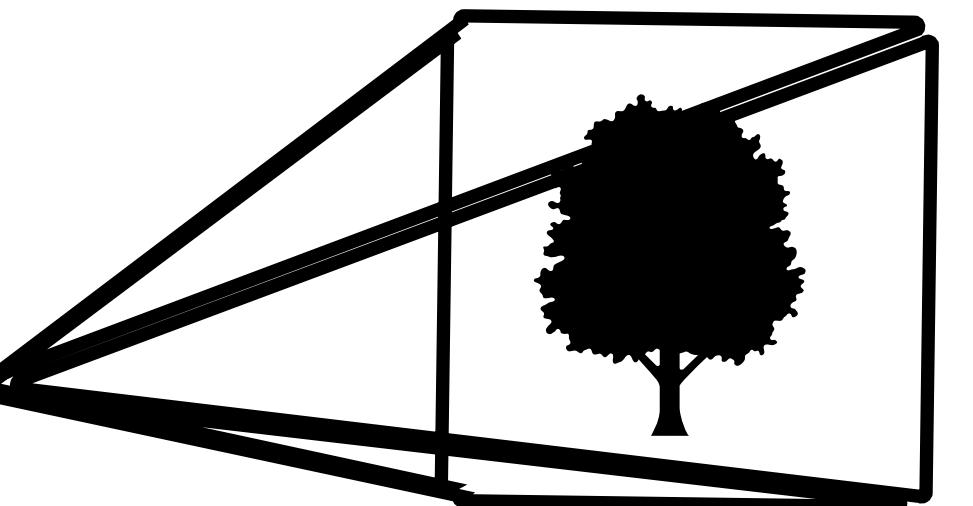
Rasterization

Isn't Path Tracing Enough?

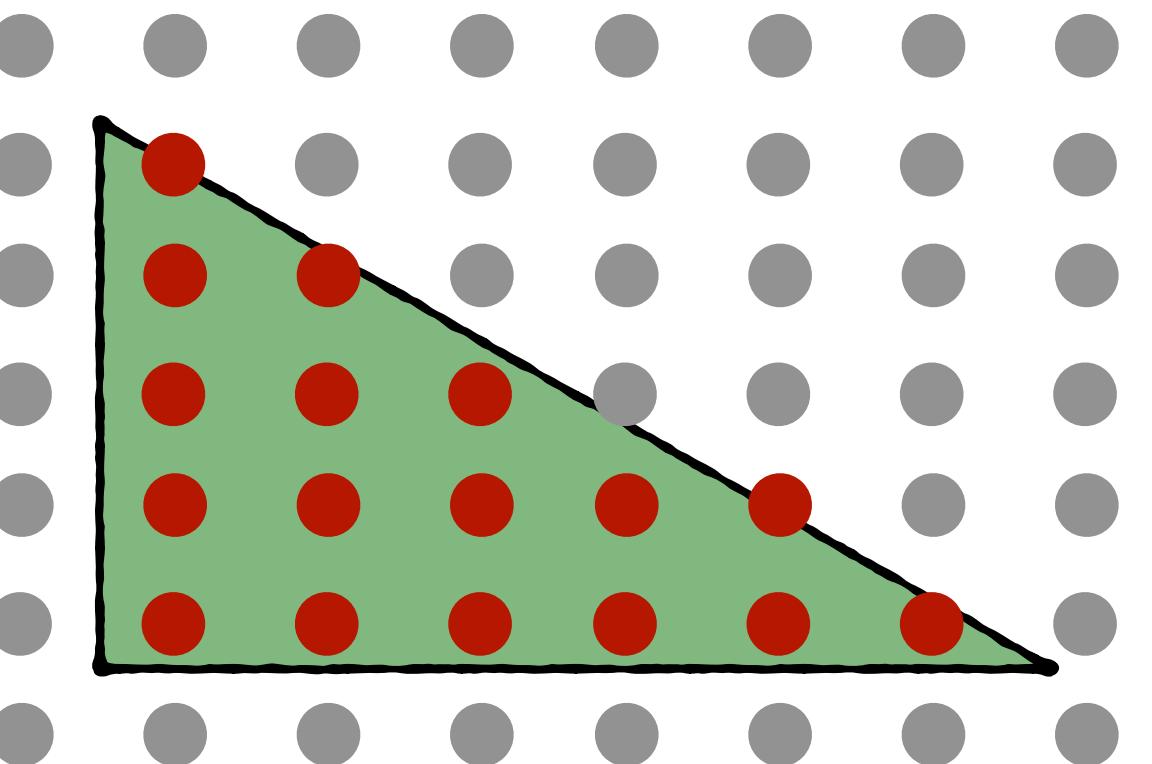
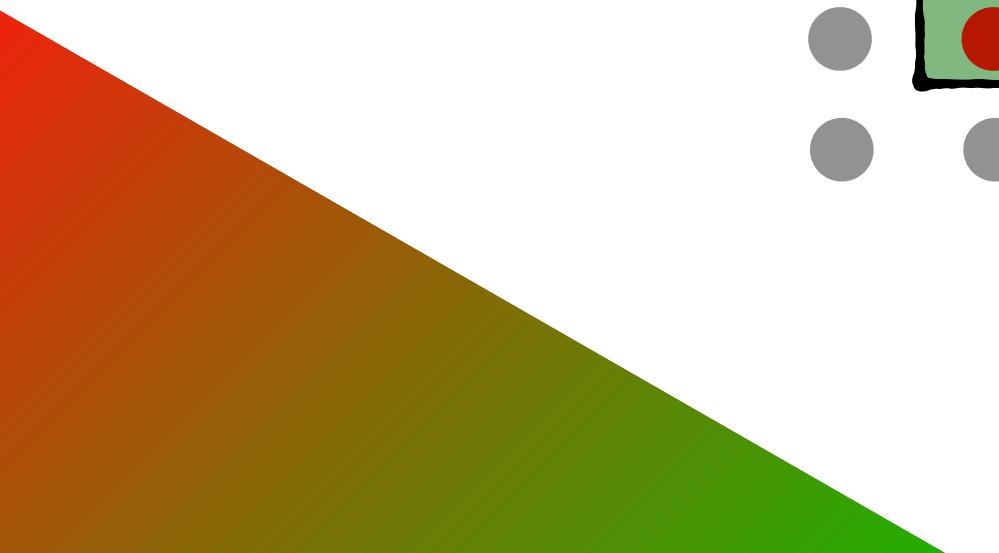
- Path tracing is slow and noisy
 - Predominantly used in cartoons & movies
 - Real-time rendering is still an ongoing effort
- Rasterization is an alternative rendering approach
 - Hard to make photorealistic
 - The main approach behind real-time graphics



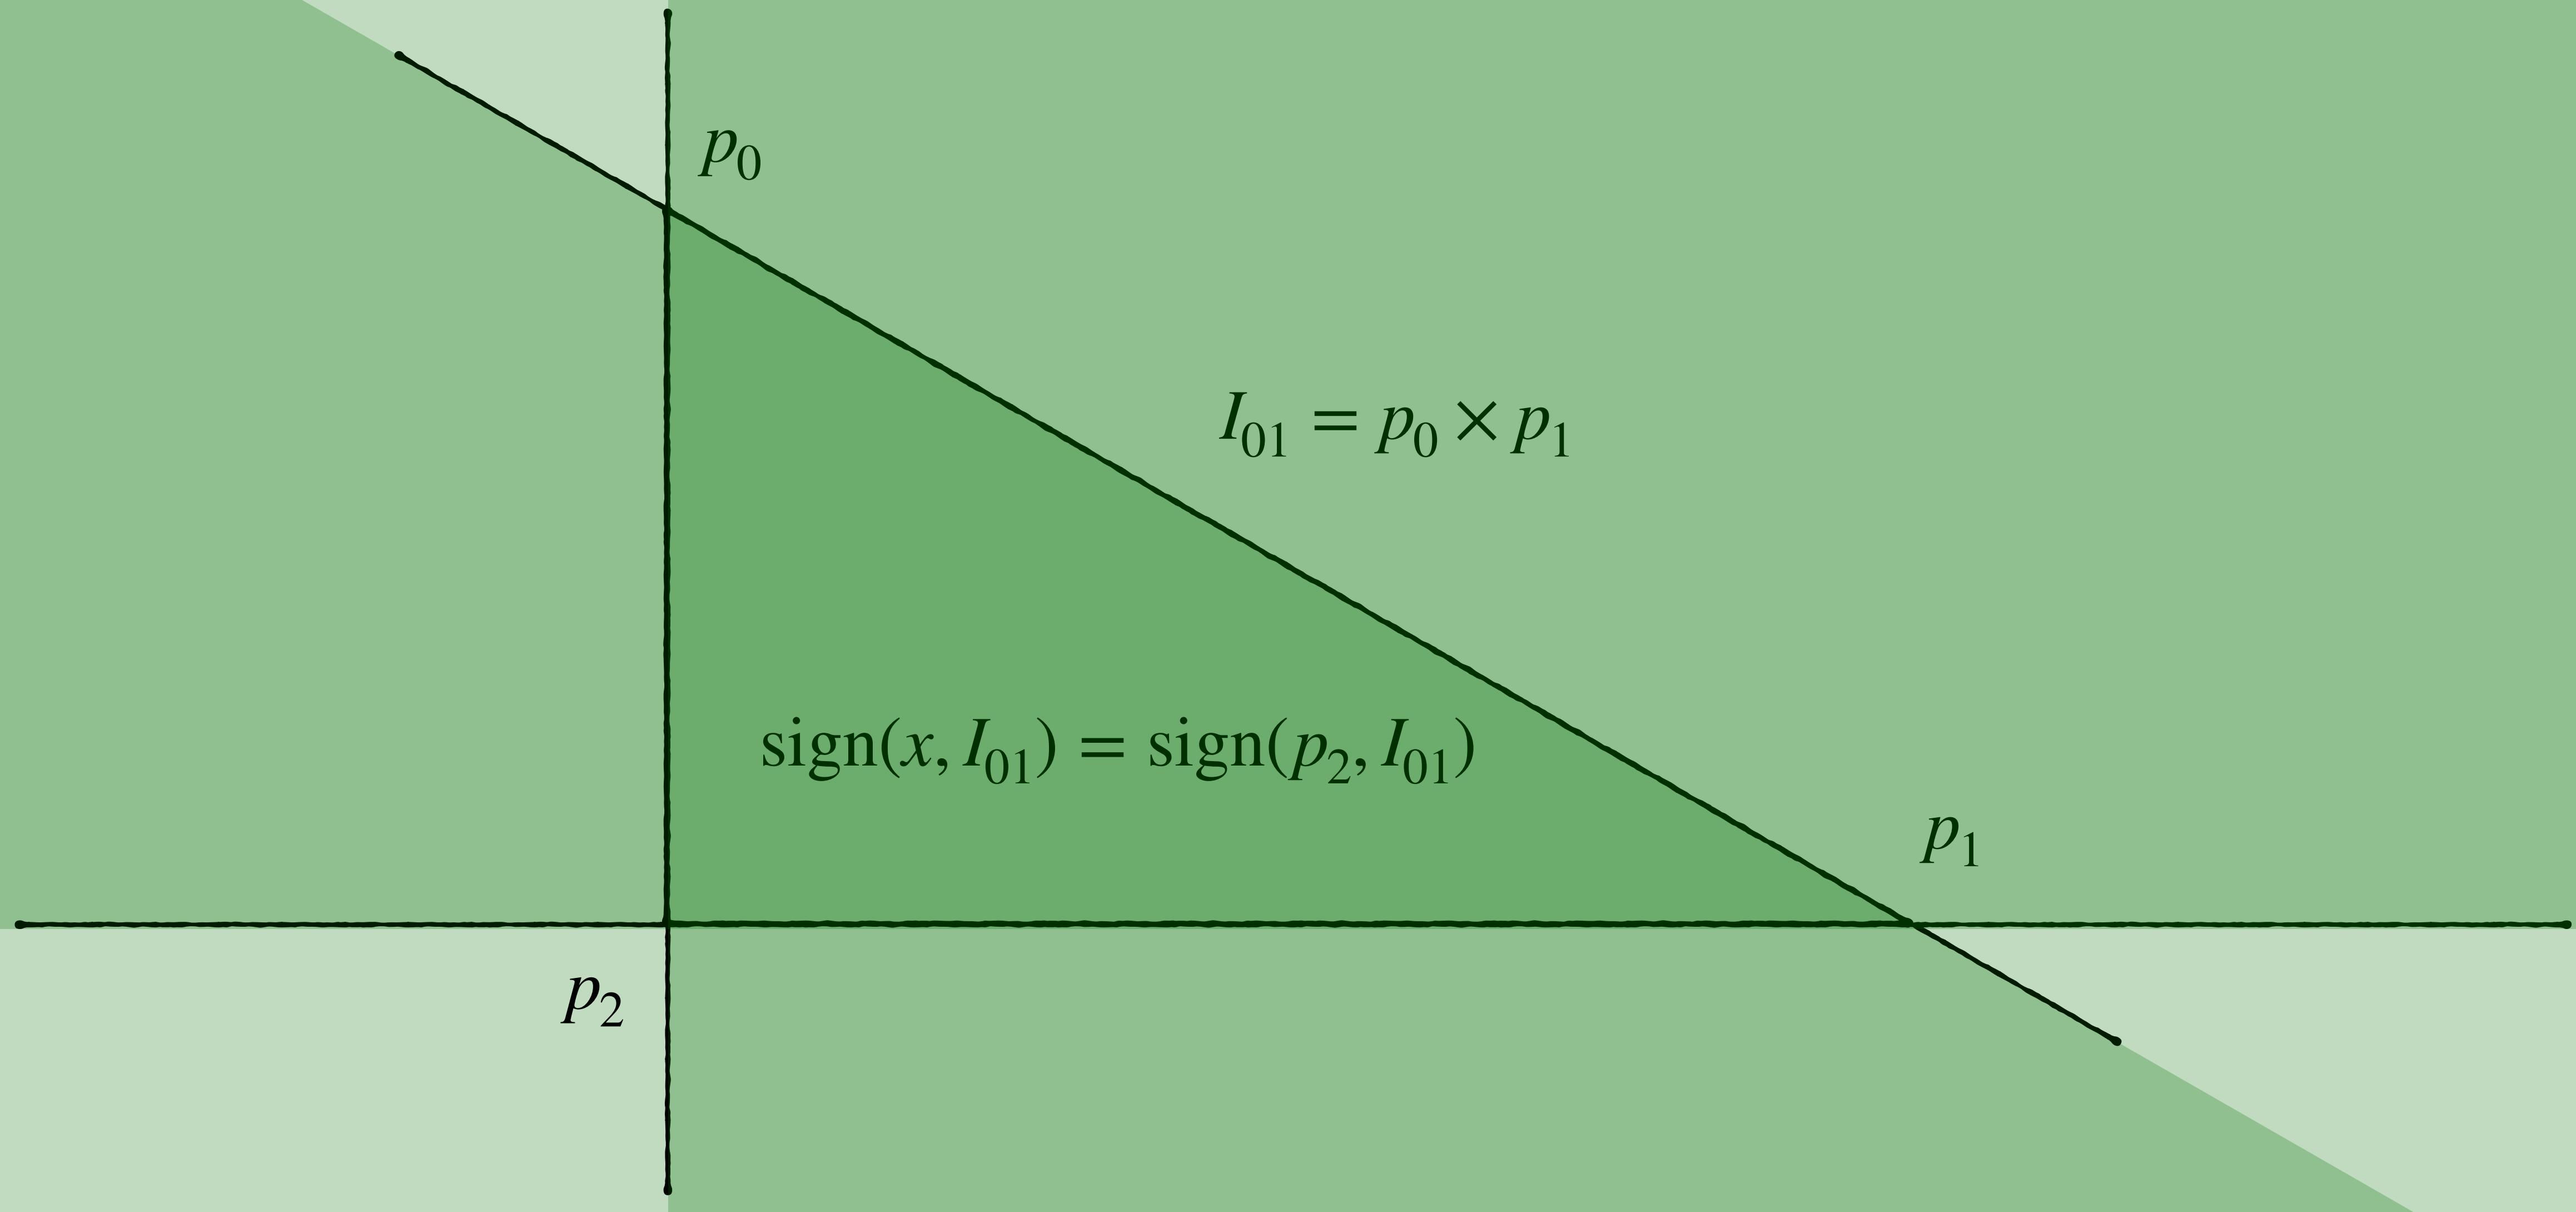
Rasterisation Pipeline



- Position objects in the world C^{w2c}
- Project objects onto the screen
- For each triangle
 - Sample coverage
 - Interpolate triangle attributes
 - Sample textures / evaluate shaders
 - Combine into the final image



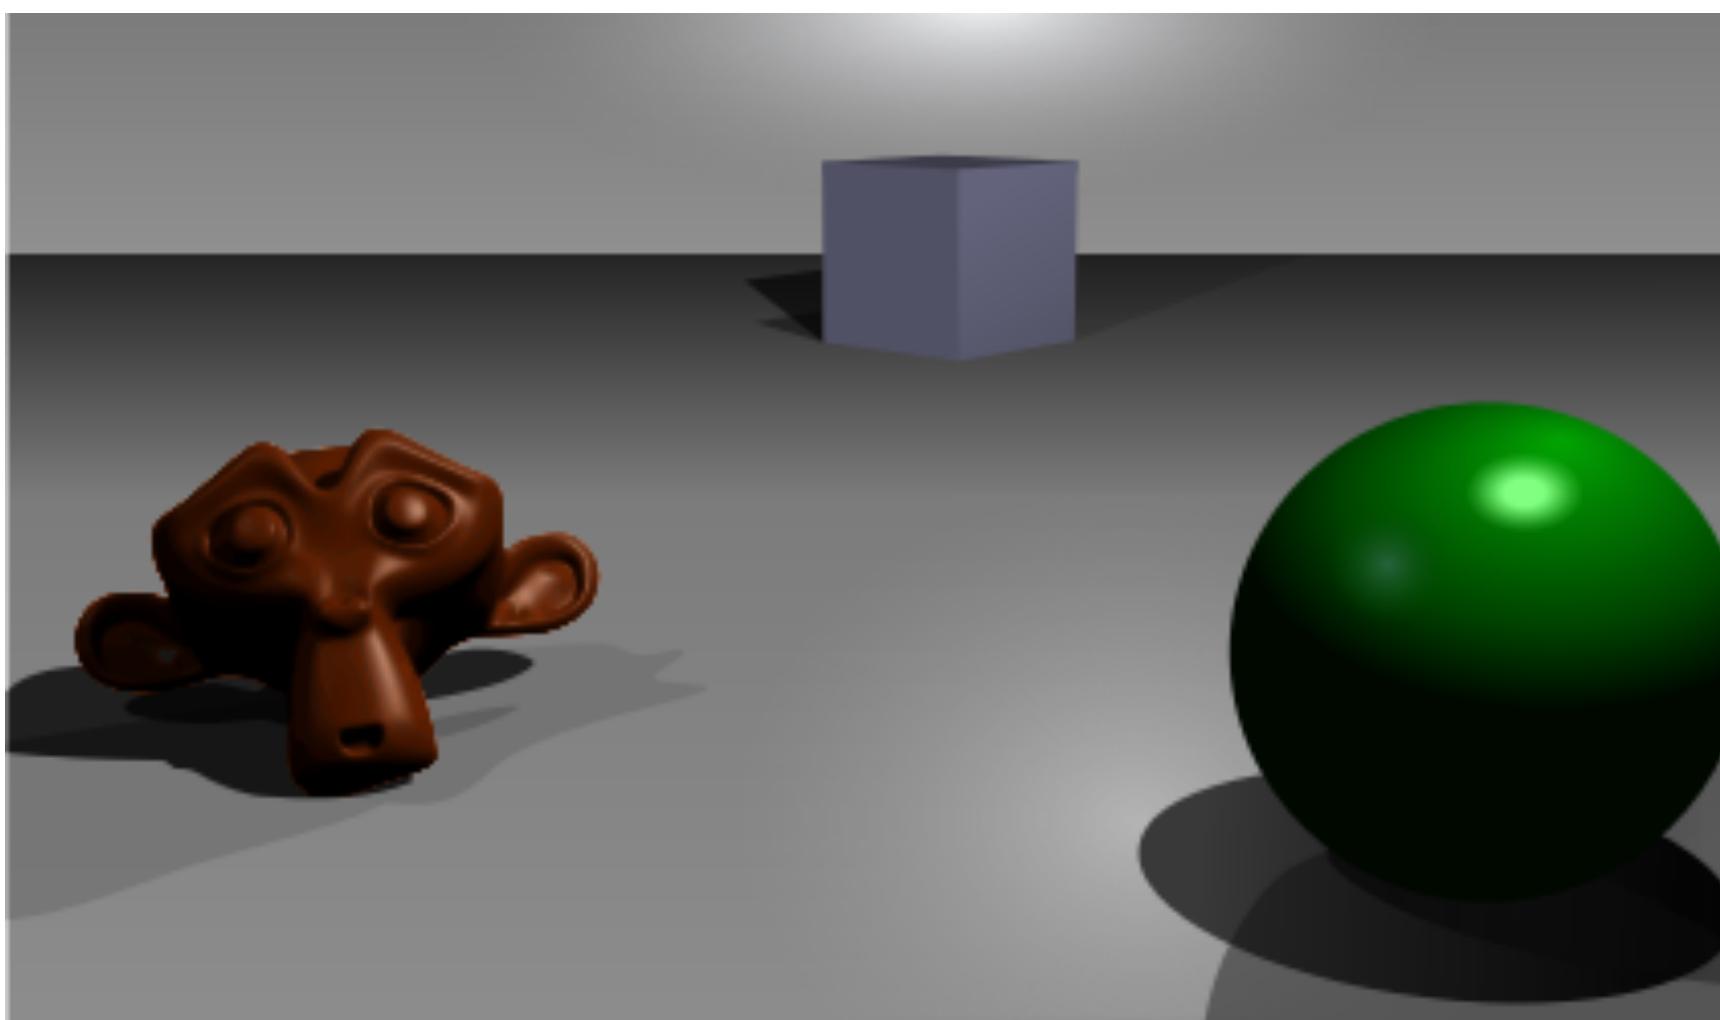
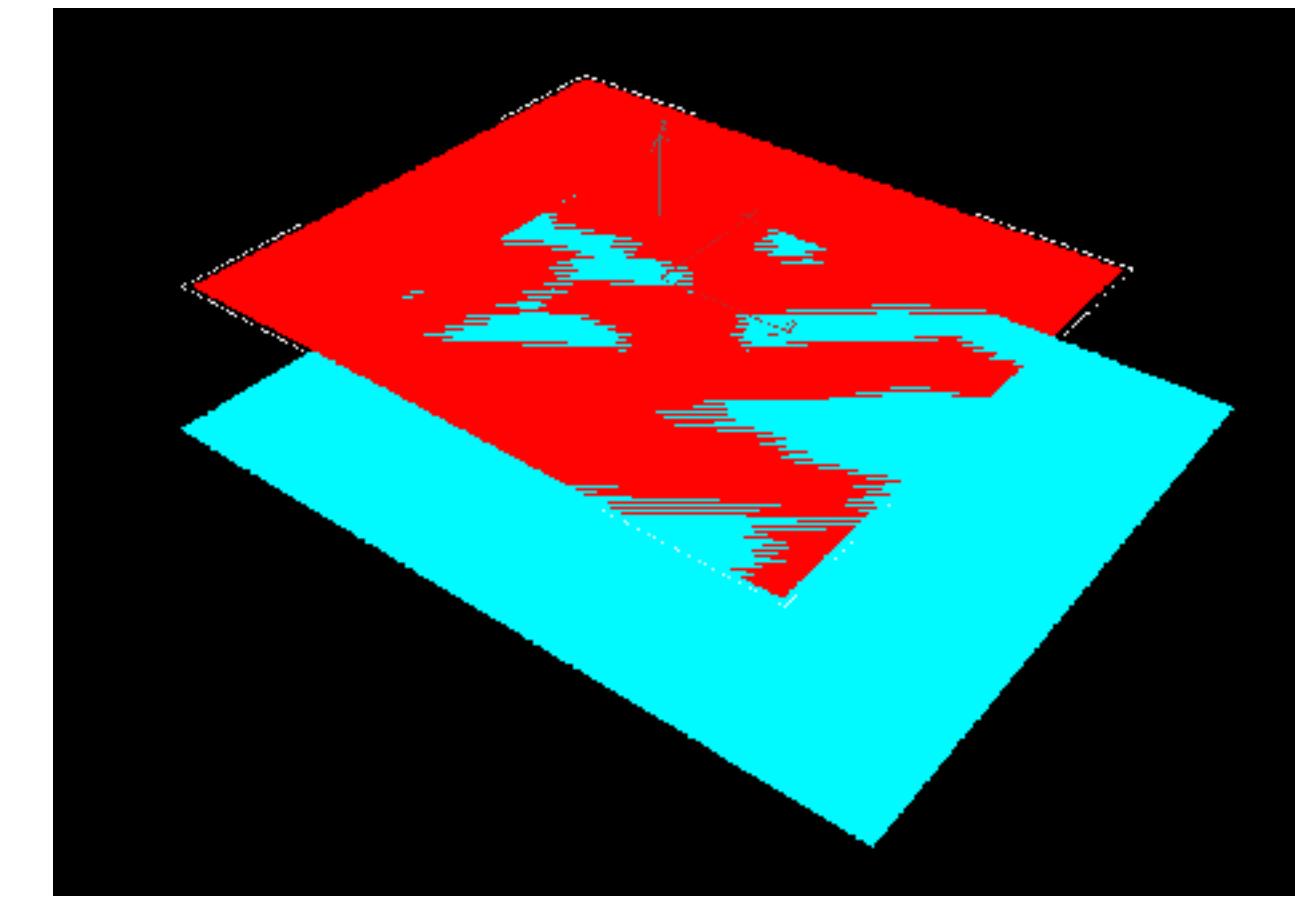
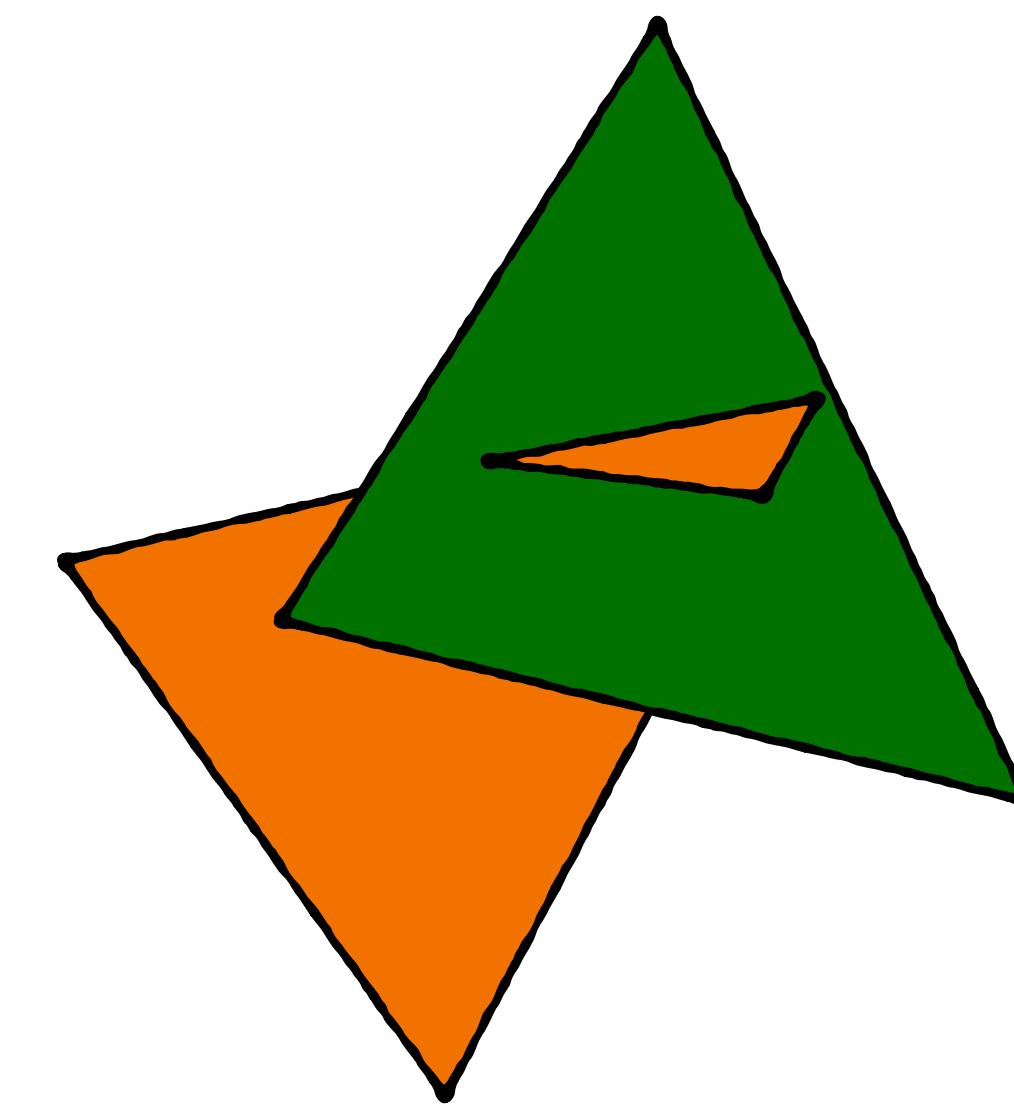
Drawing a Single Triangle



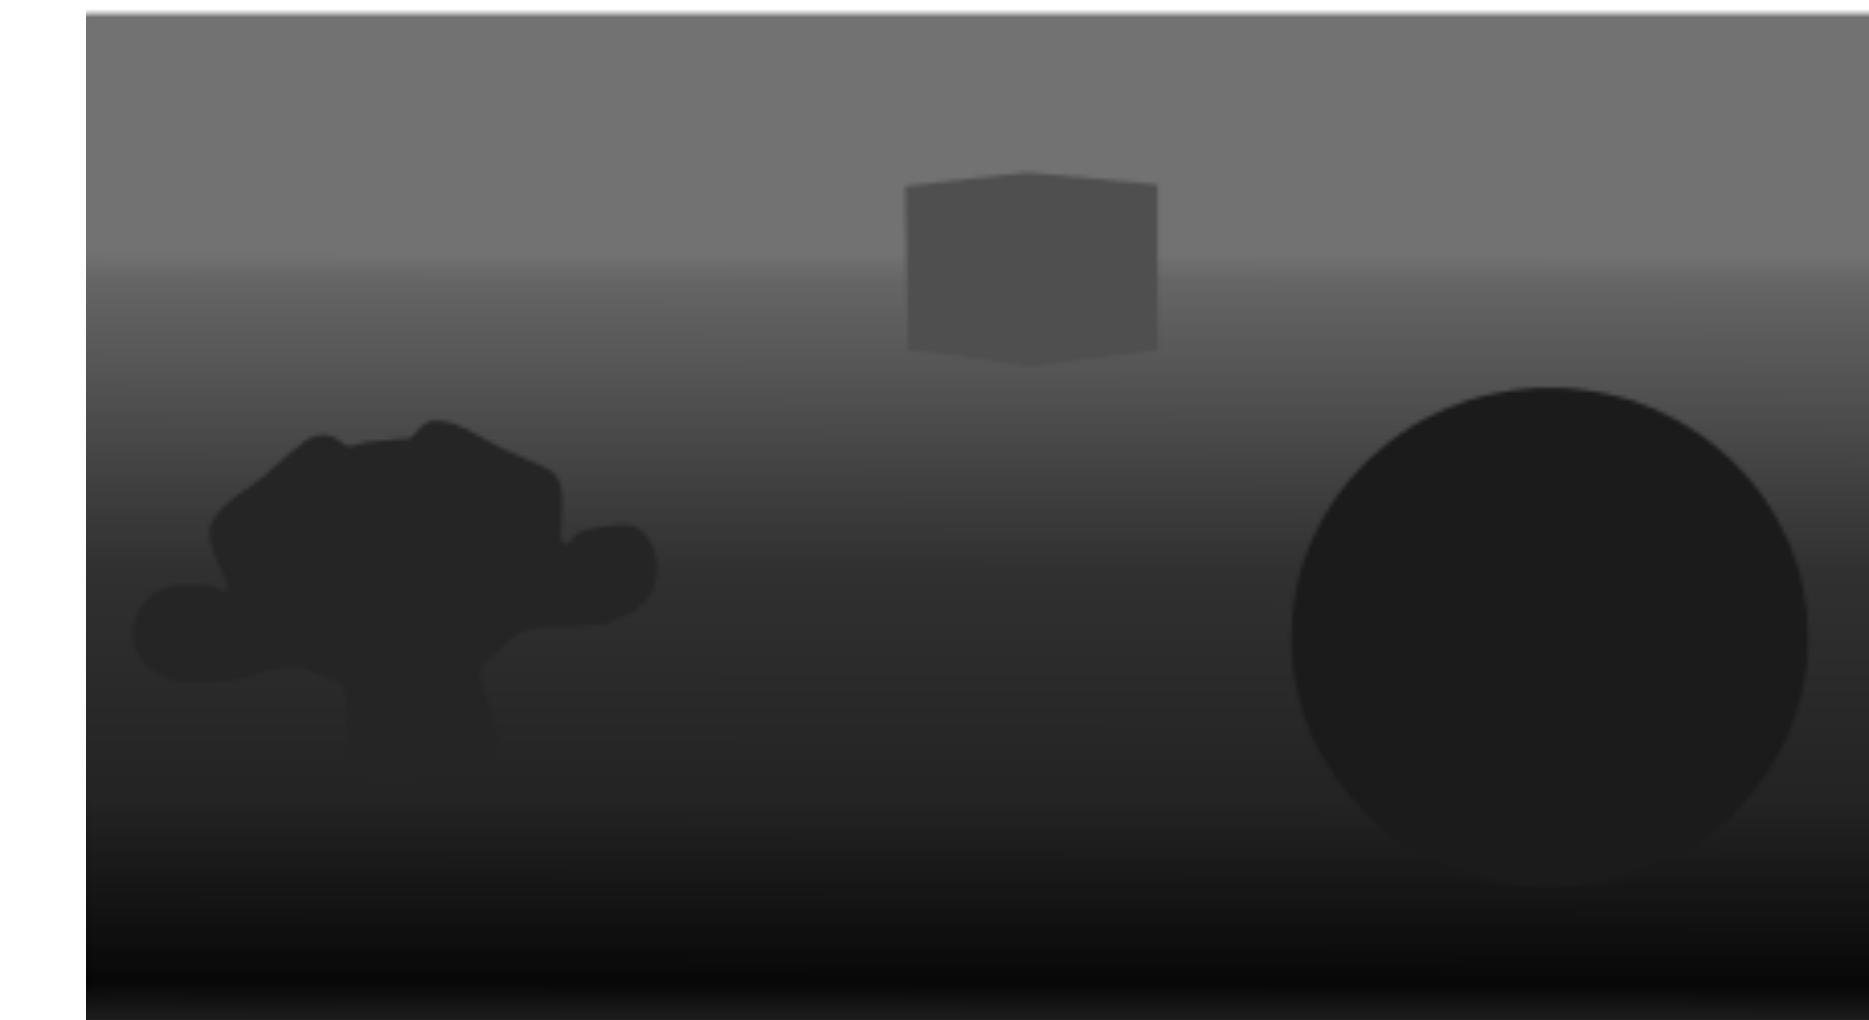
Drawing Two Triangles

Depth Buffer

- Two triangles may intersect
- Need to render only the closest points
- Solution: store depth buffer



A simple three-dimensional scene



Z-buffer representation

Extended Projection Matrix

- Projection matrices we encountered omitted one coordinate

$$\begin{pmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$$

- The omitted coordinate stores depth
- Modify matrix to store depth

$$\begin{pmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & \frac{z_n + z_f}{z_n - z_f} & \frac{2z_f z_n}{z_n - z_f} \\ 0 & 0 & -1 & 0 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$$

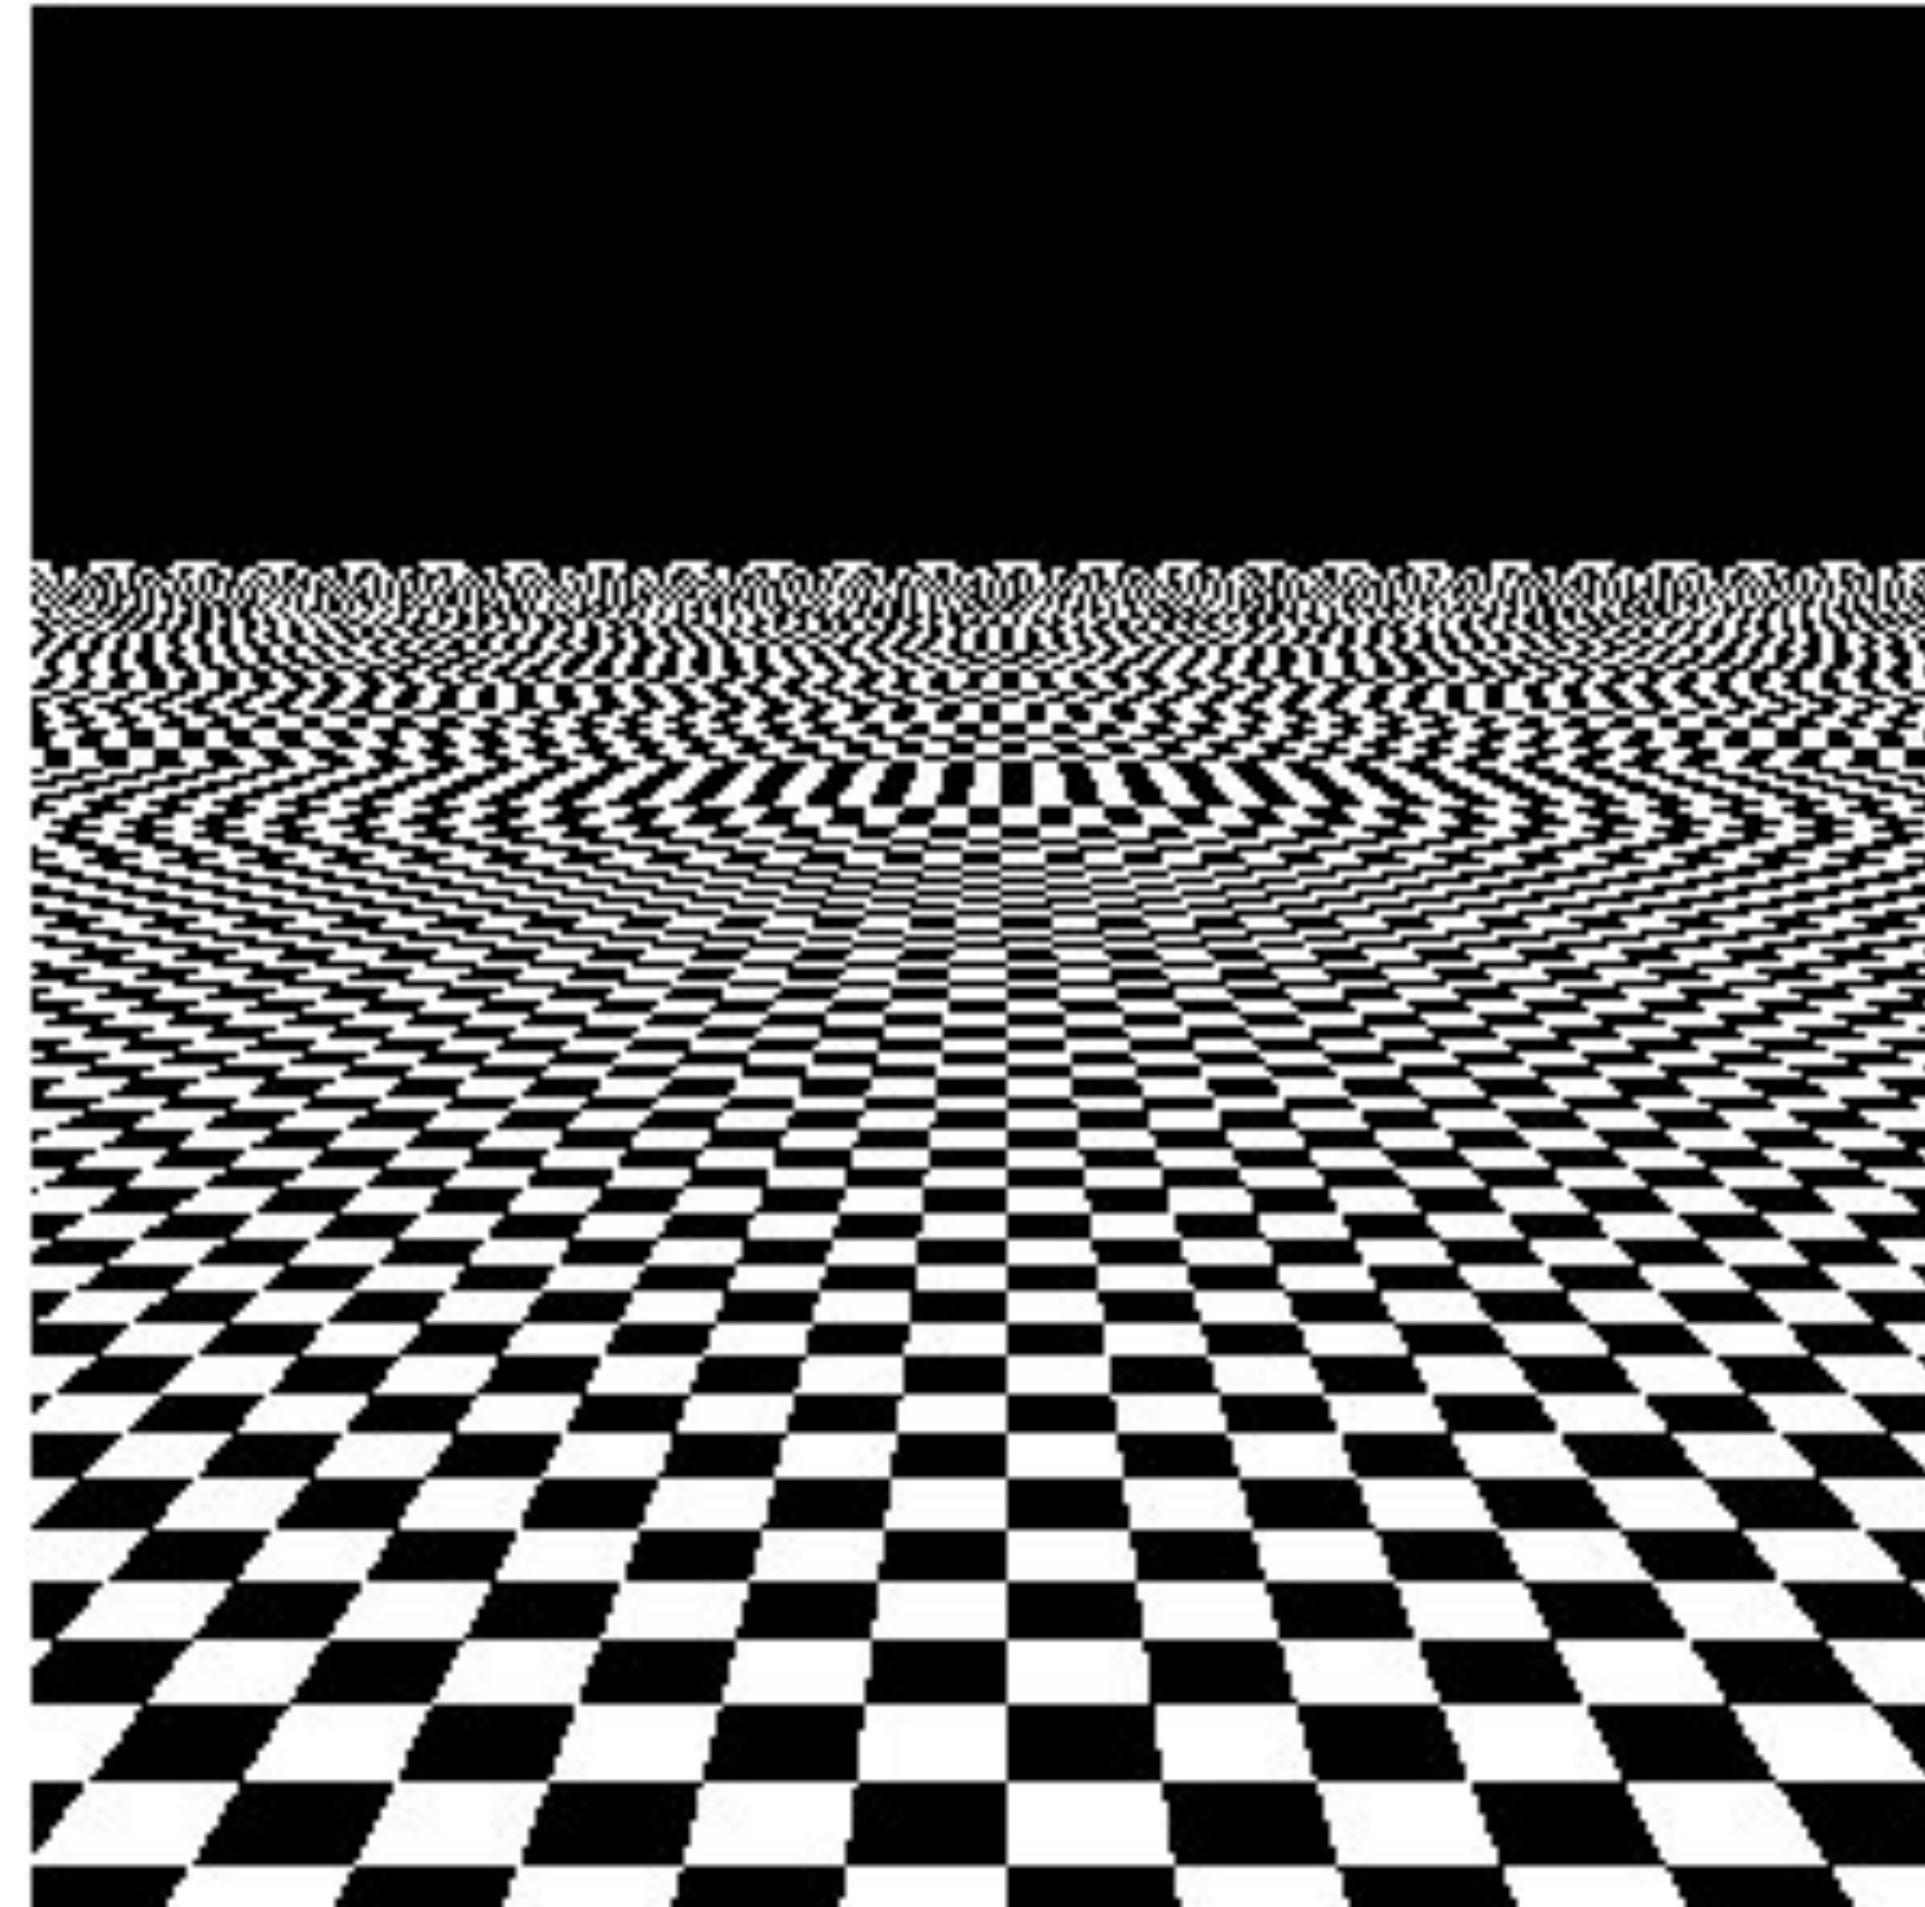
- Maps depth $[z_n, z_f]$ into $[-1, 1]$

Pseudo-Code for Rasterisation

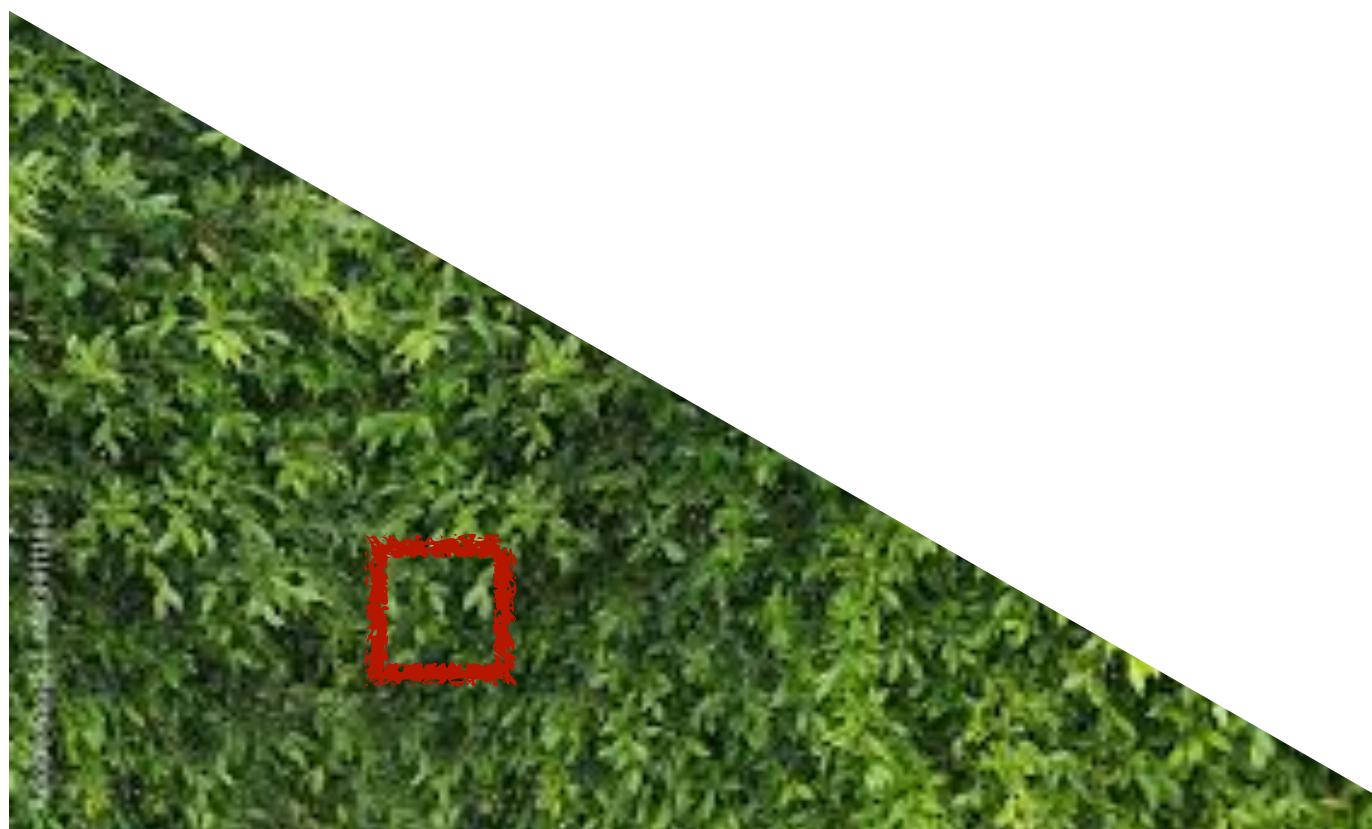
```
for each triangle O in scene # from closest to farthest
    project and transform O onto the screen
    for each pixel (x, y)
        if (x, y) is in O and (x, y) is closest so far
            compute color at (x, y)
            save corresponding depth to depth buffer
```

Aliasing

One of the Many Details Swept Under the Carpet



Colour of a Point



pixel #1



pixel #2

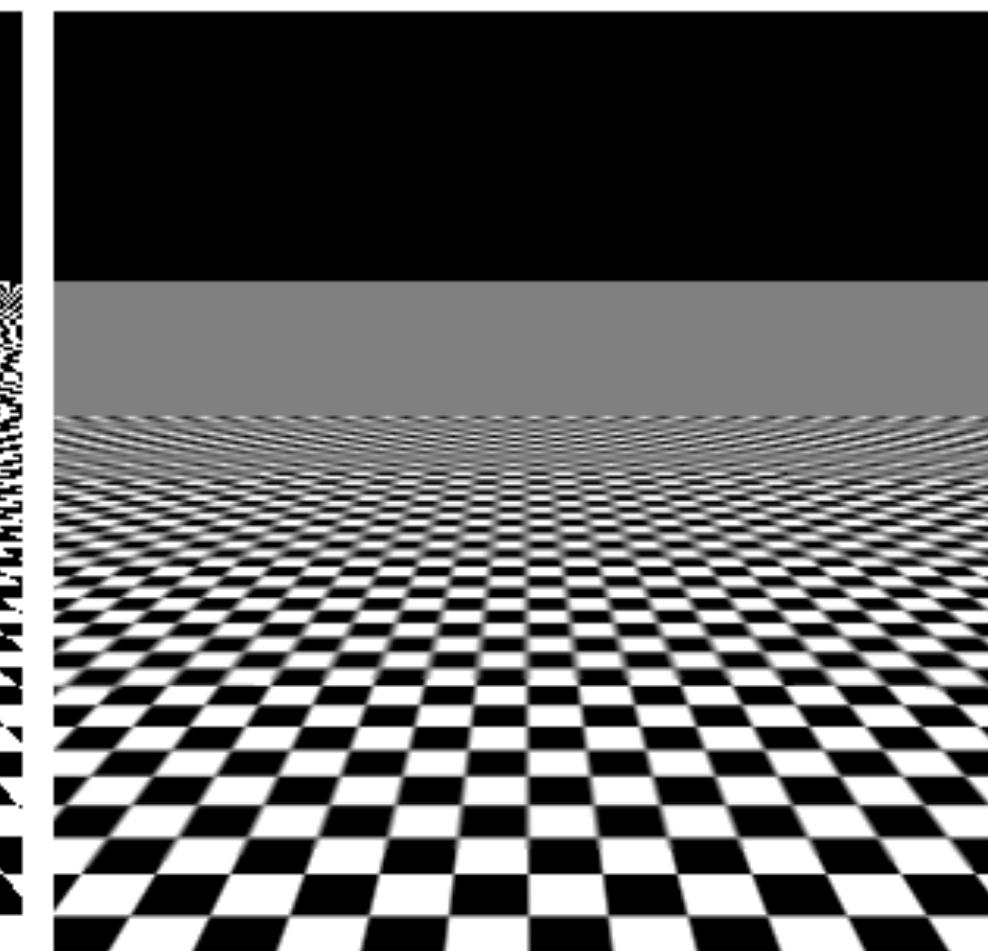
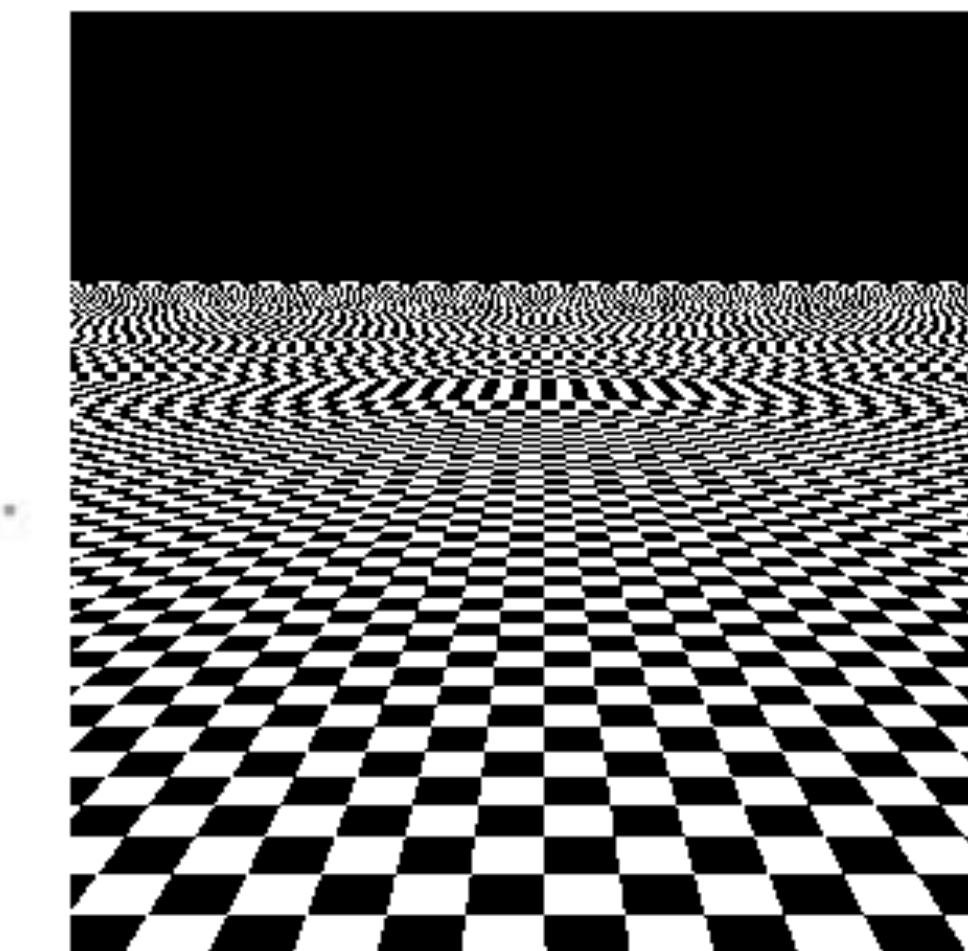
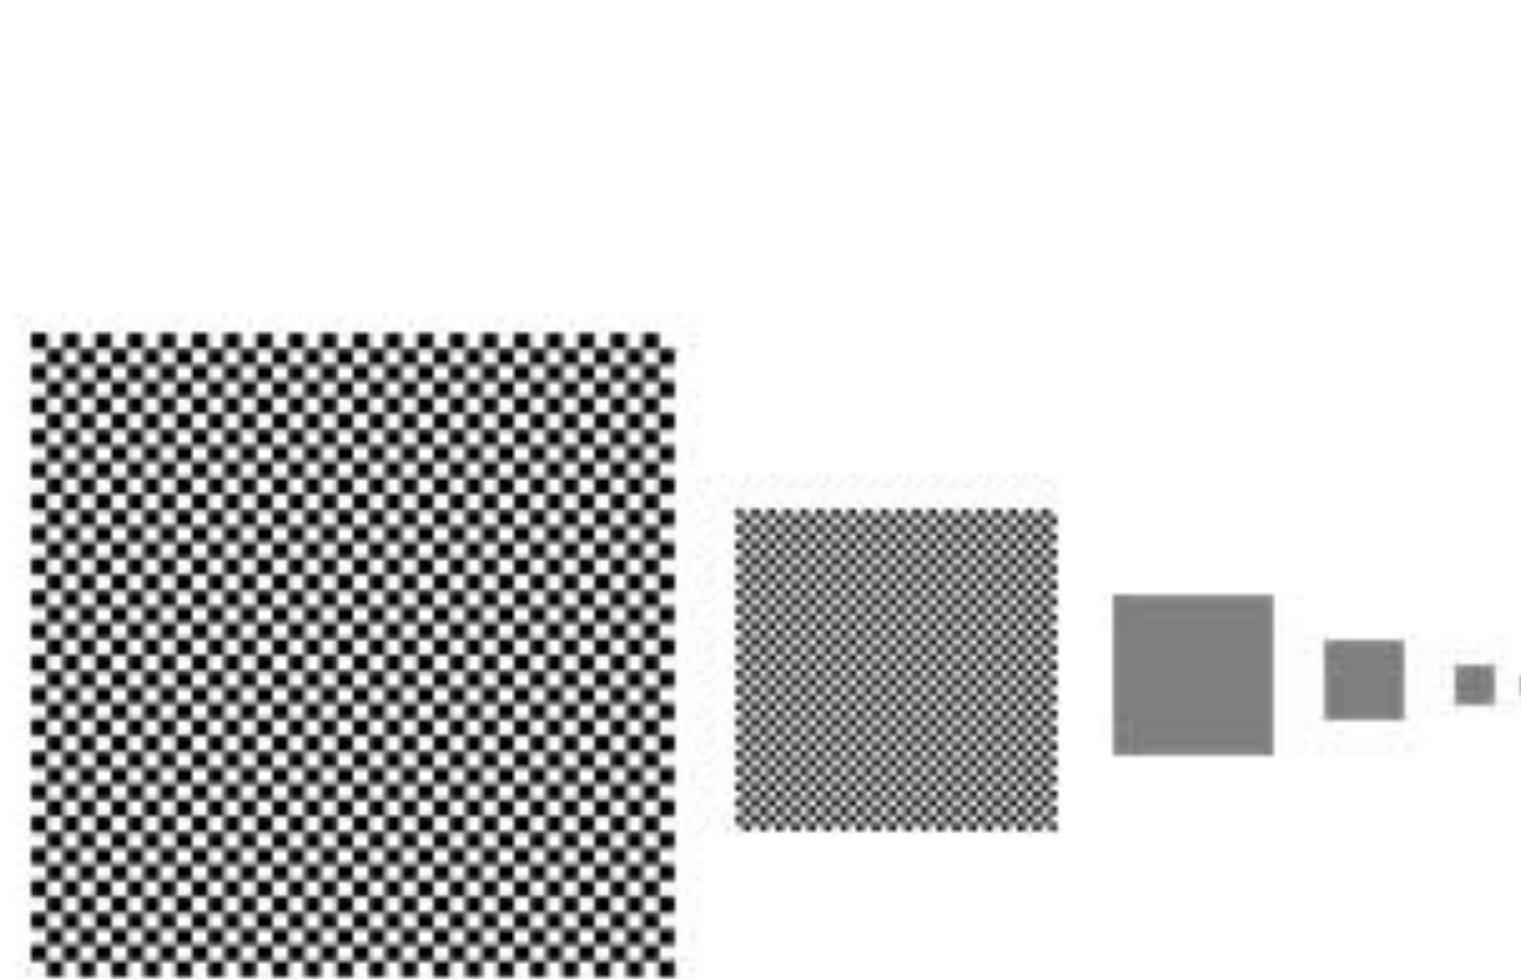
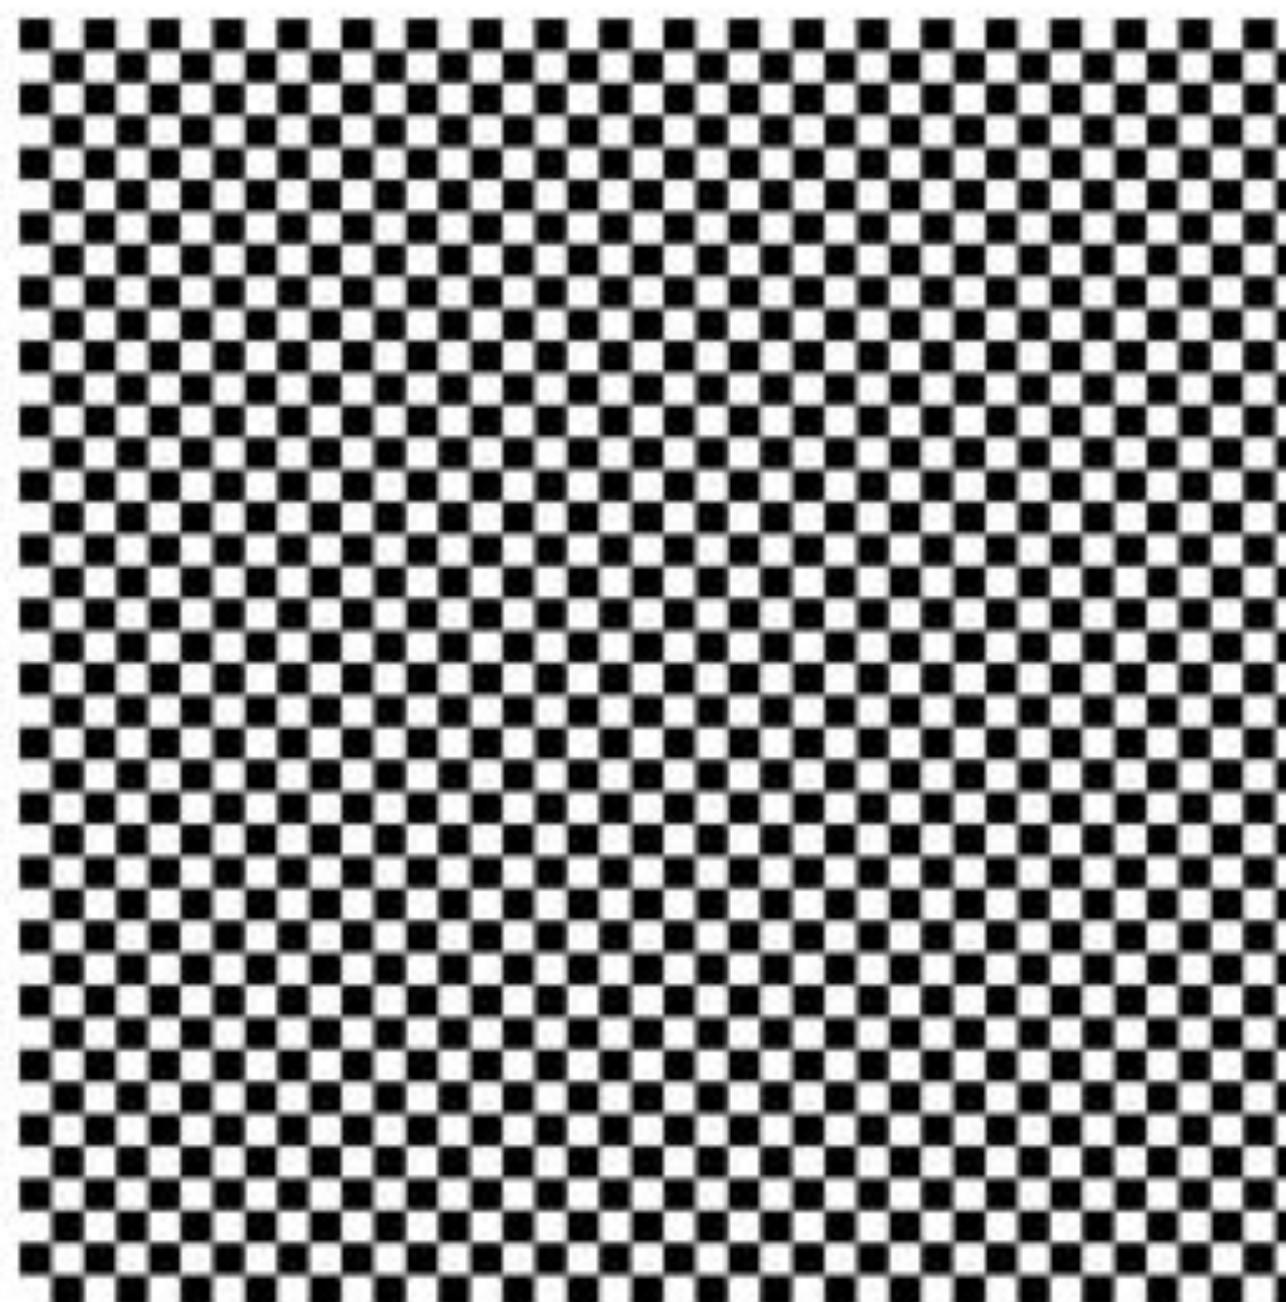


pixel #3

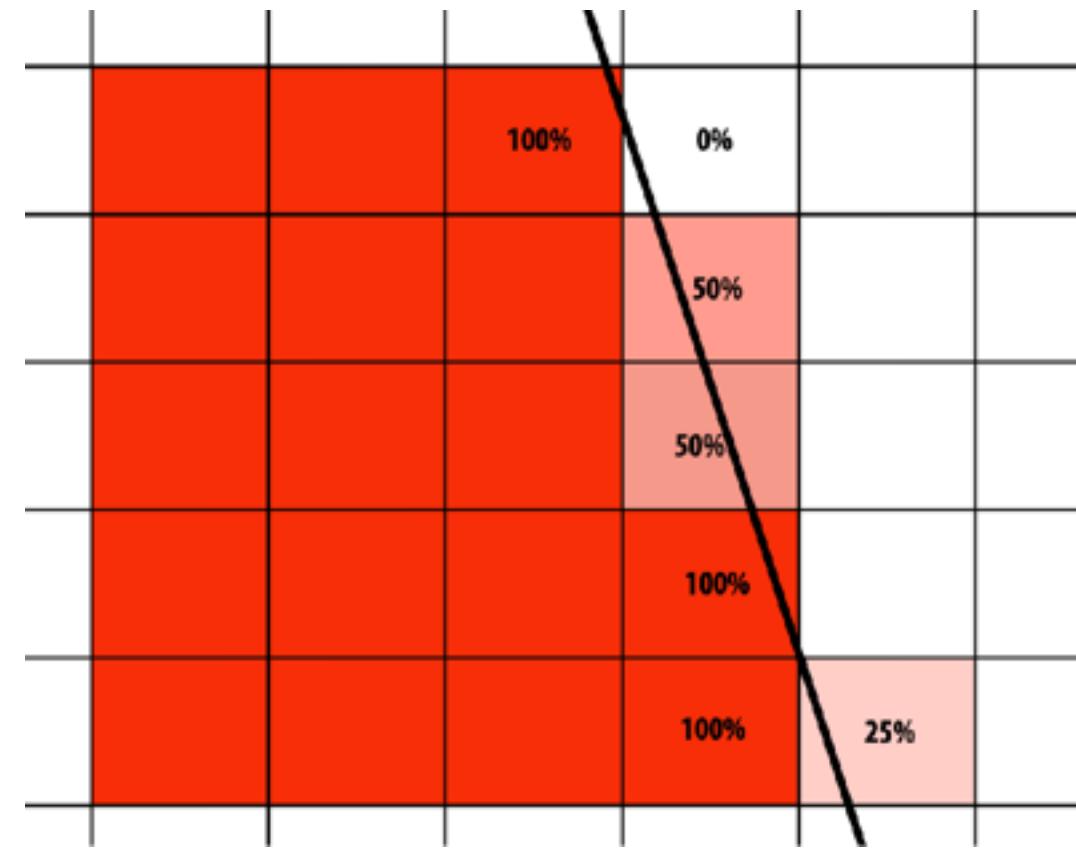
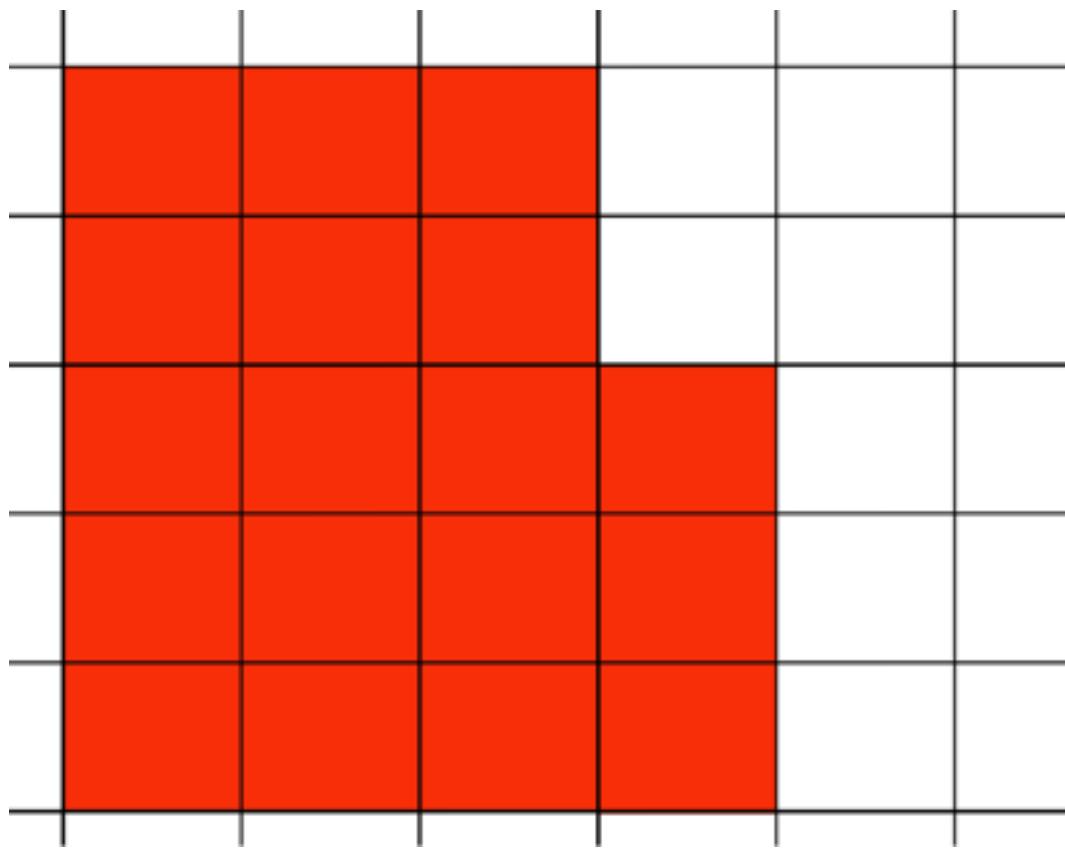
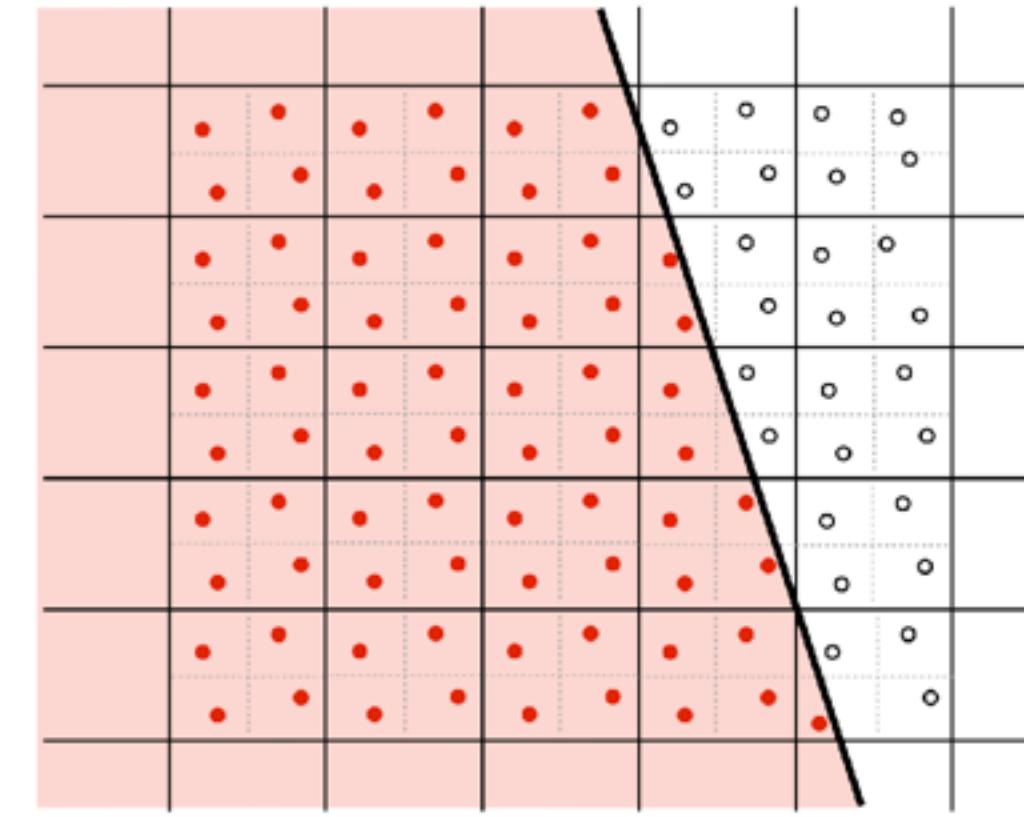
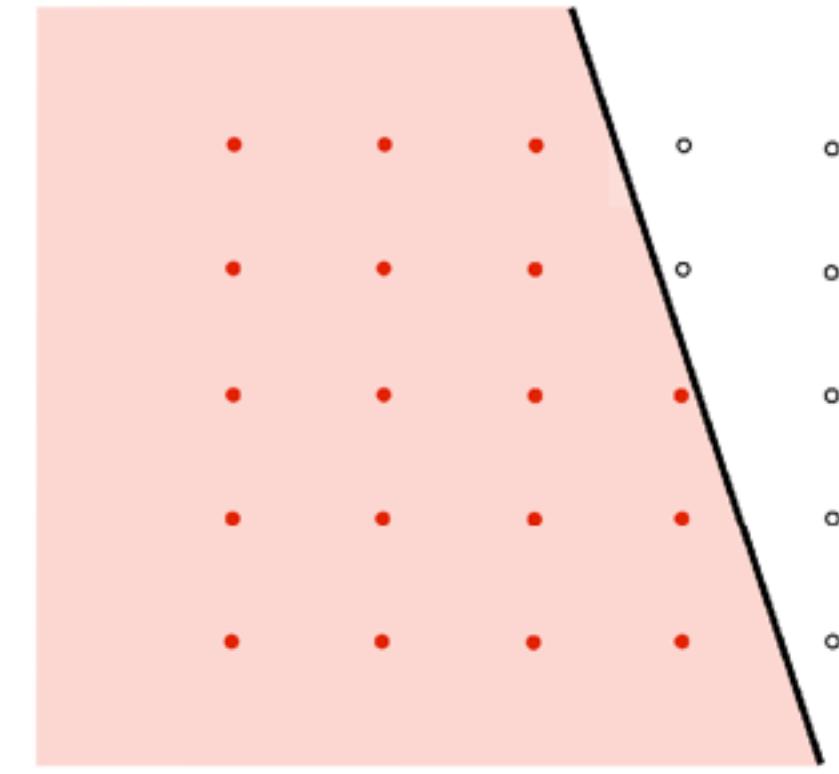
- Ideally, the resulting color should be the average color across the pixel

MIP-Mapping

- Computing the average for each pixel would be too expensive
- Solution: pre-compute texture at different scales



Border Aliasing



nvdiffrast

Vertex processing



Rasterization



Interpolation



Texturing



Shading



Anti-Aliasing



Applications: nvdiffrec



Reference photo



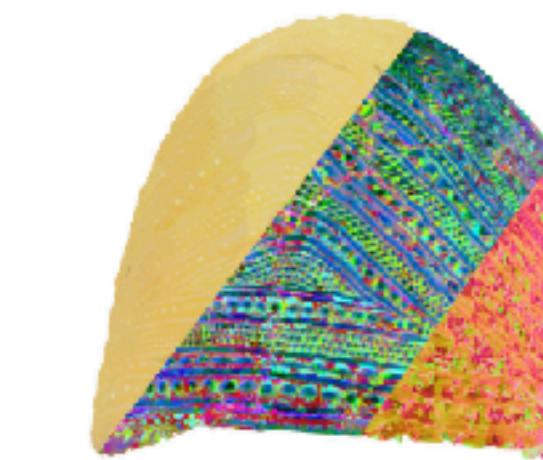
Scene editing



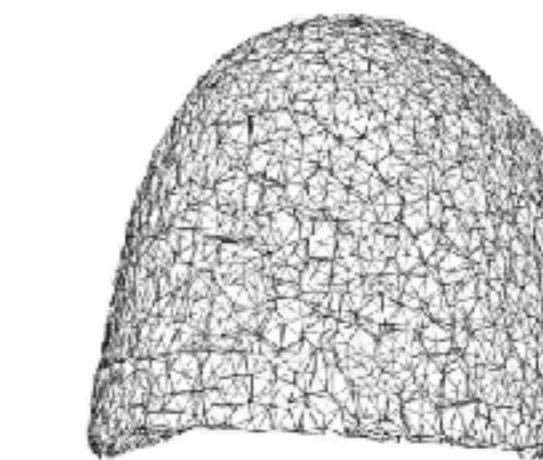
Cloth simulation



Material editing

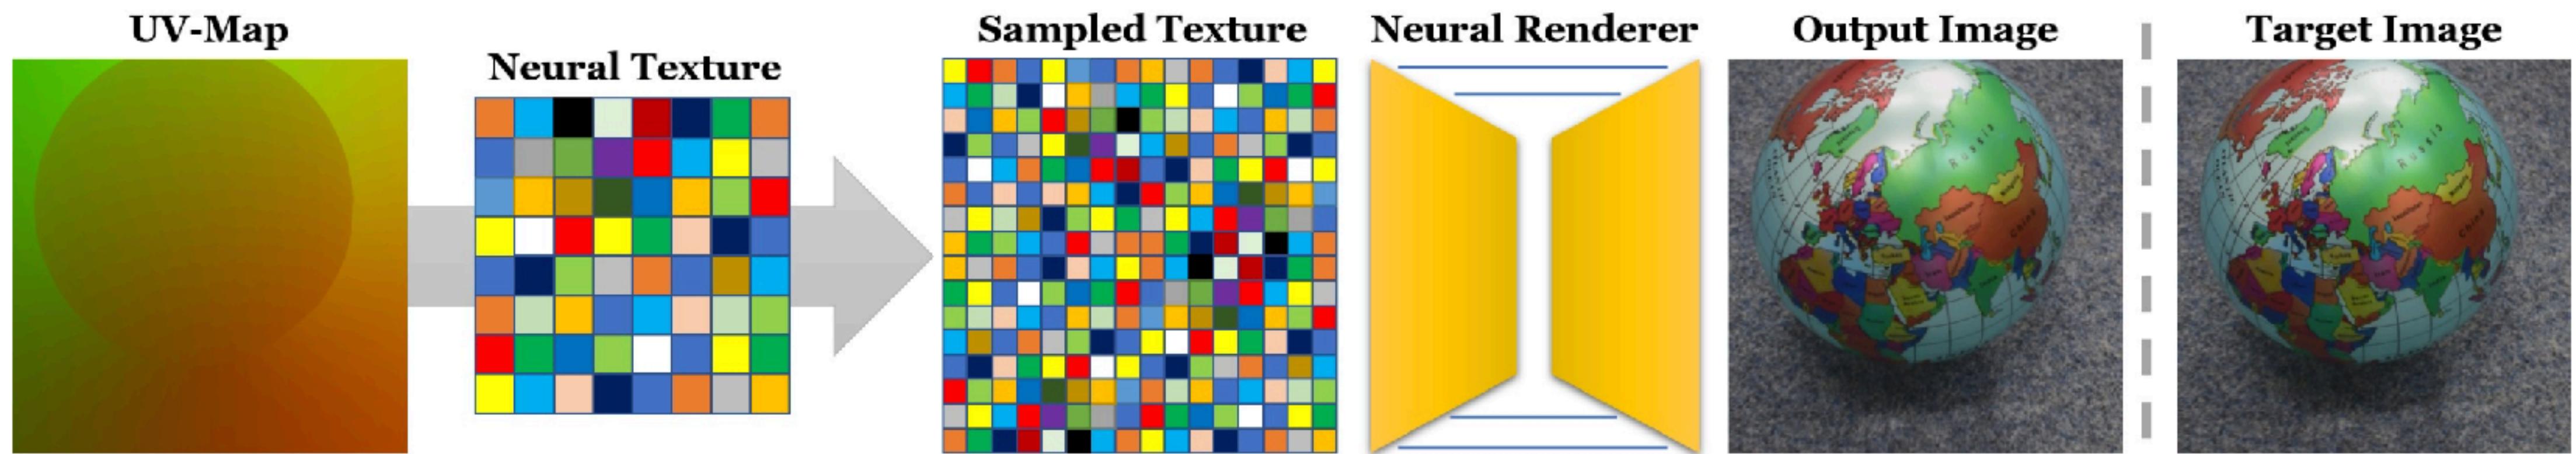


$k_d/k_{orm}/\mathbf{n}$

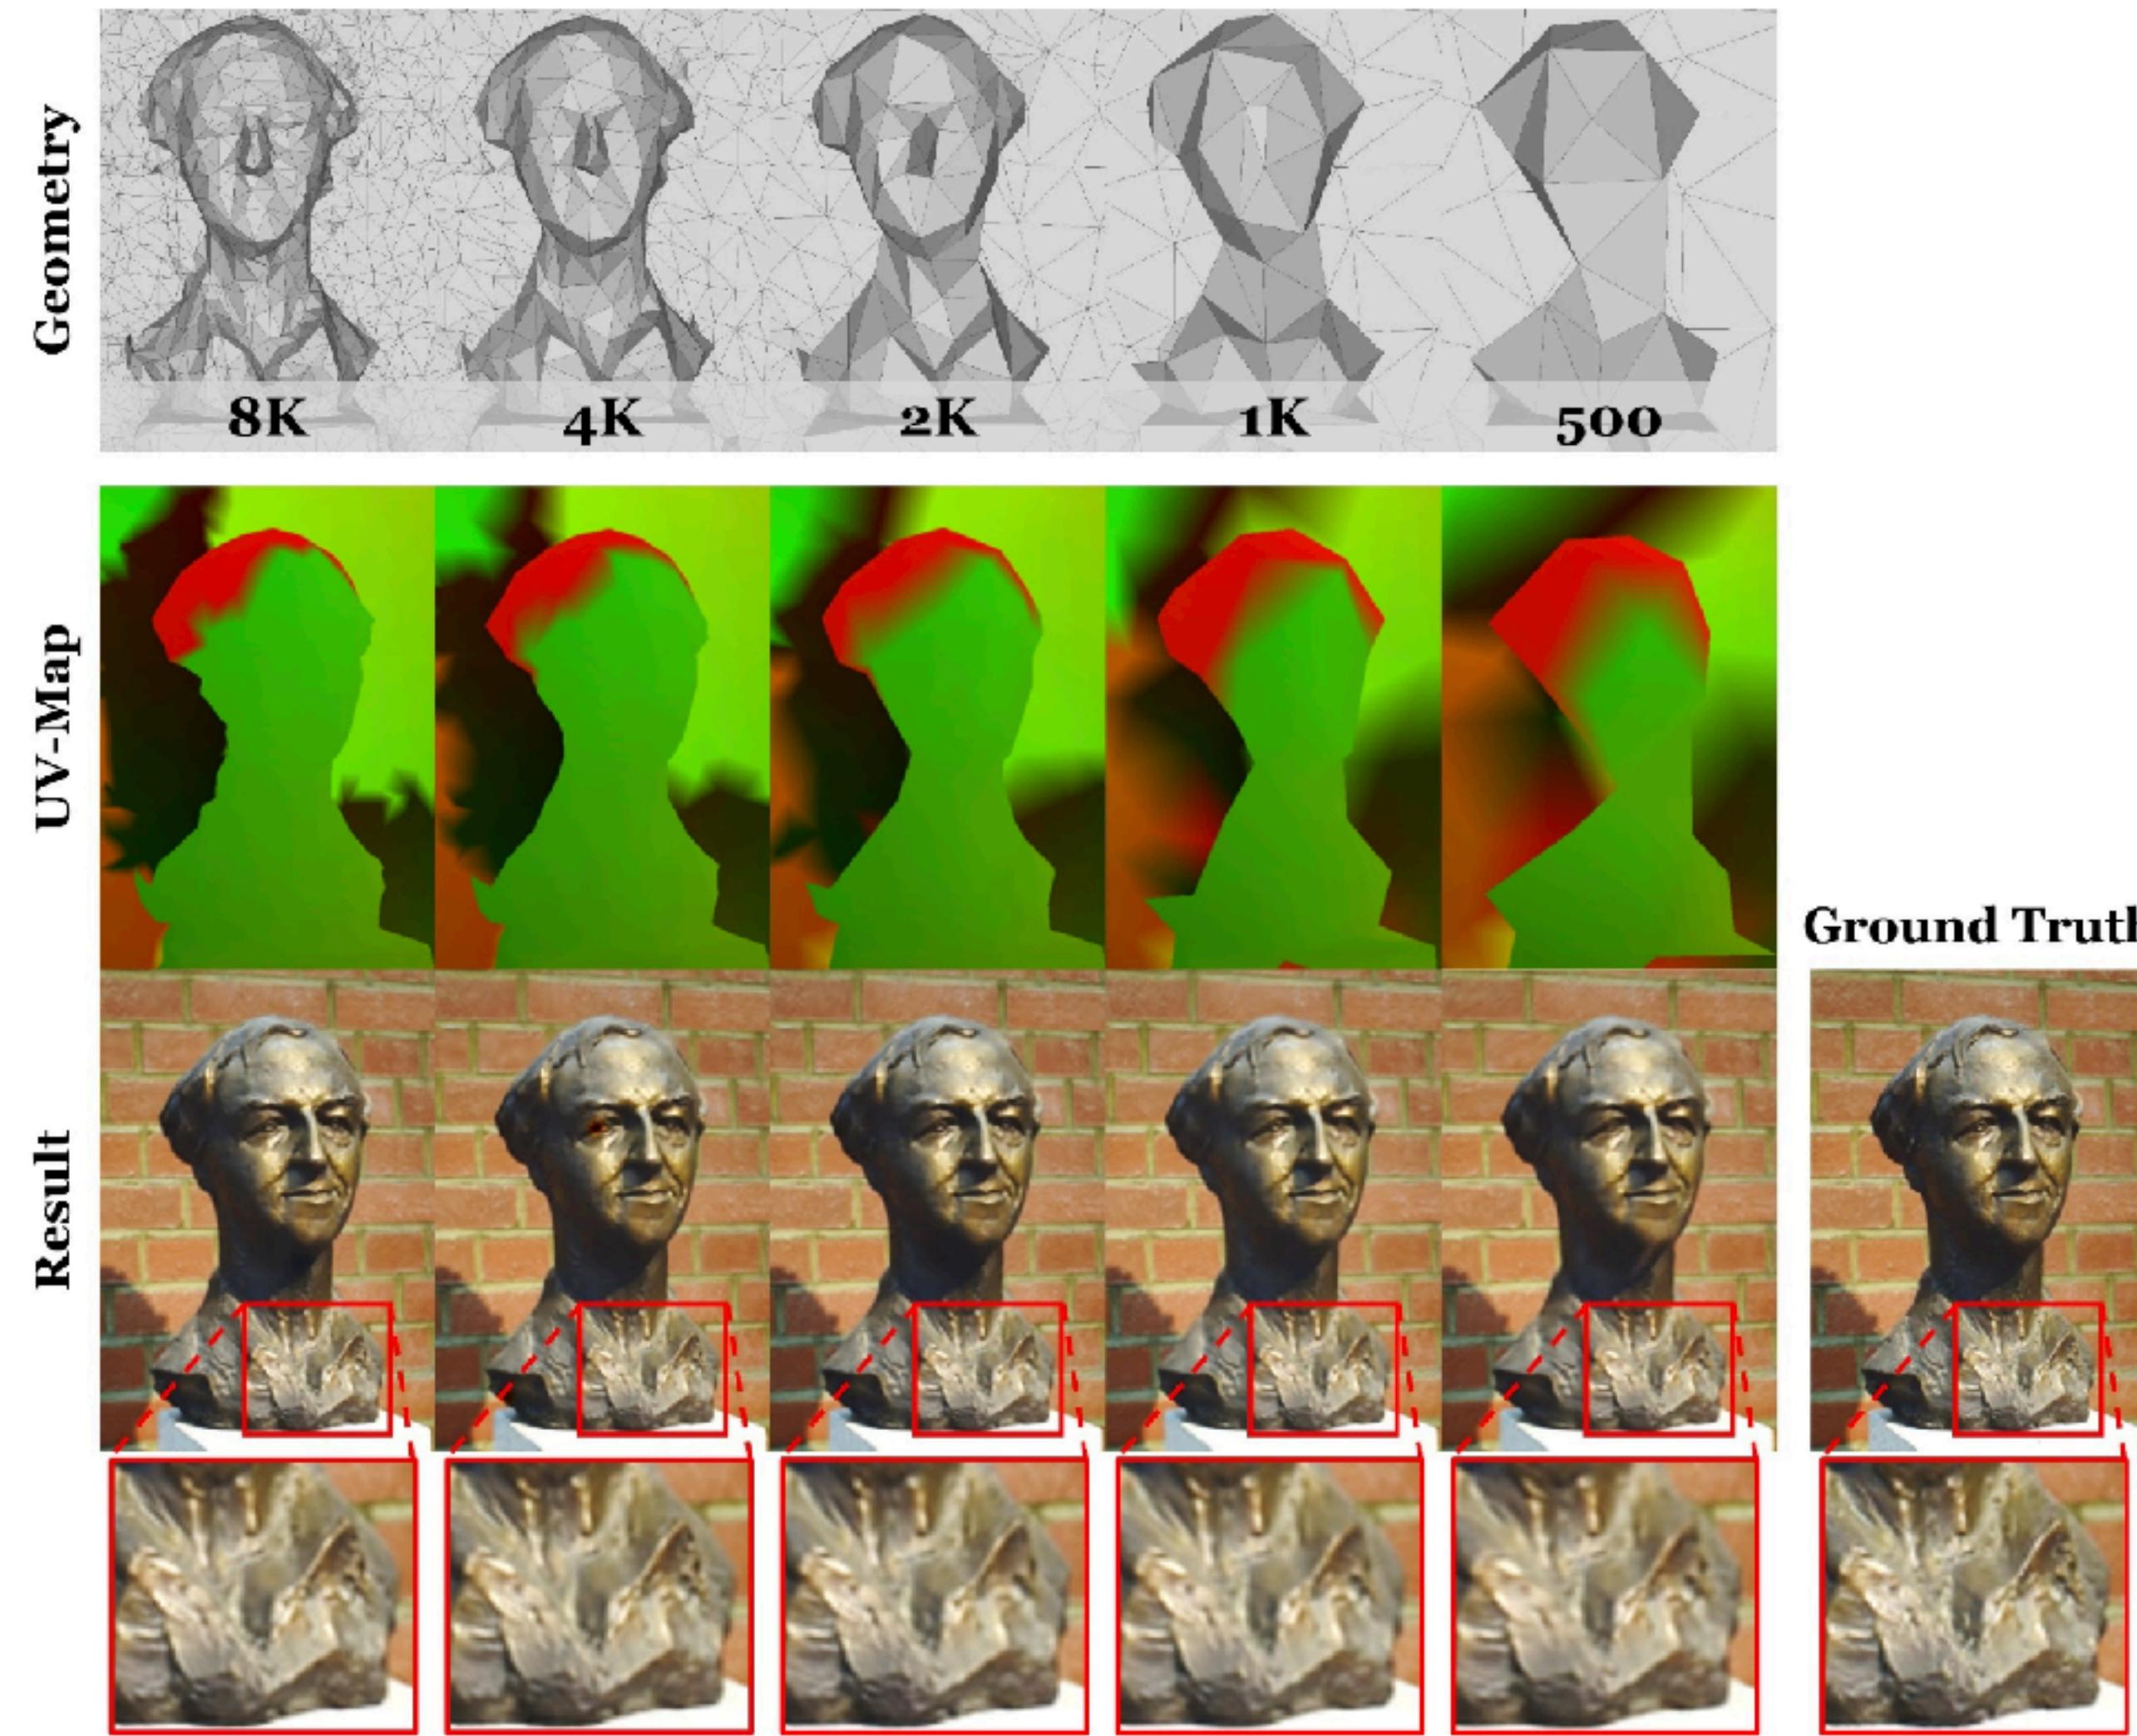


Mesh

Applications: Deferred Neural Rendering



Applications: Deferred Neural Rendering



Key Takeaways

- Path tracing
 - Photorealistic & Elegant
 - Slow
- Rasterisation
 - Hard to make realistic renders
 - Fast
- Both allow differentiation for inverse graphics, rasterisation is a little more common
- Next time: implicit representations for 3D data

